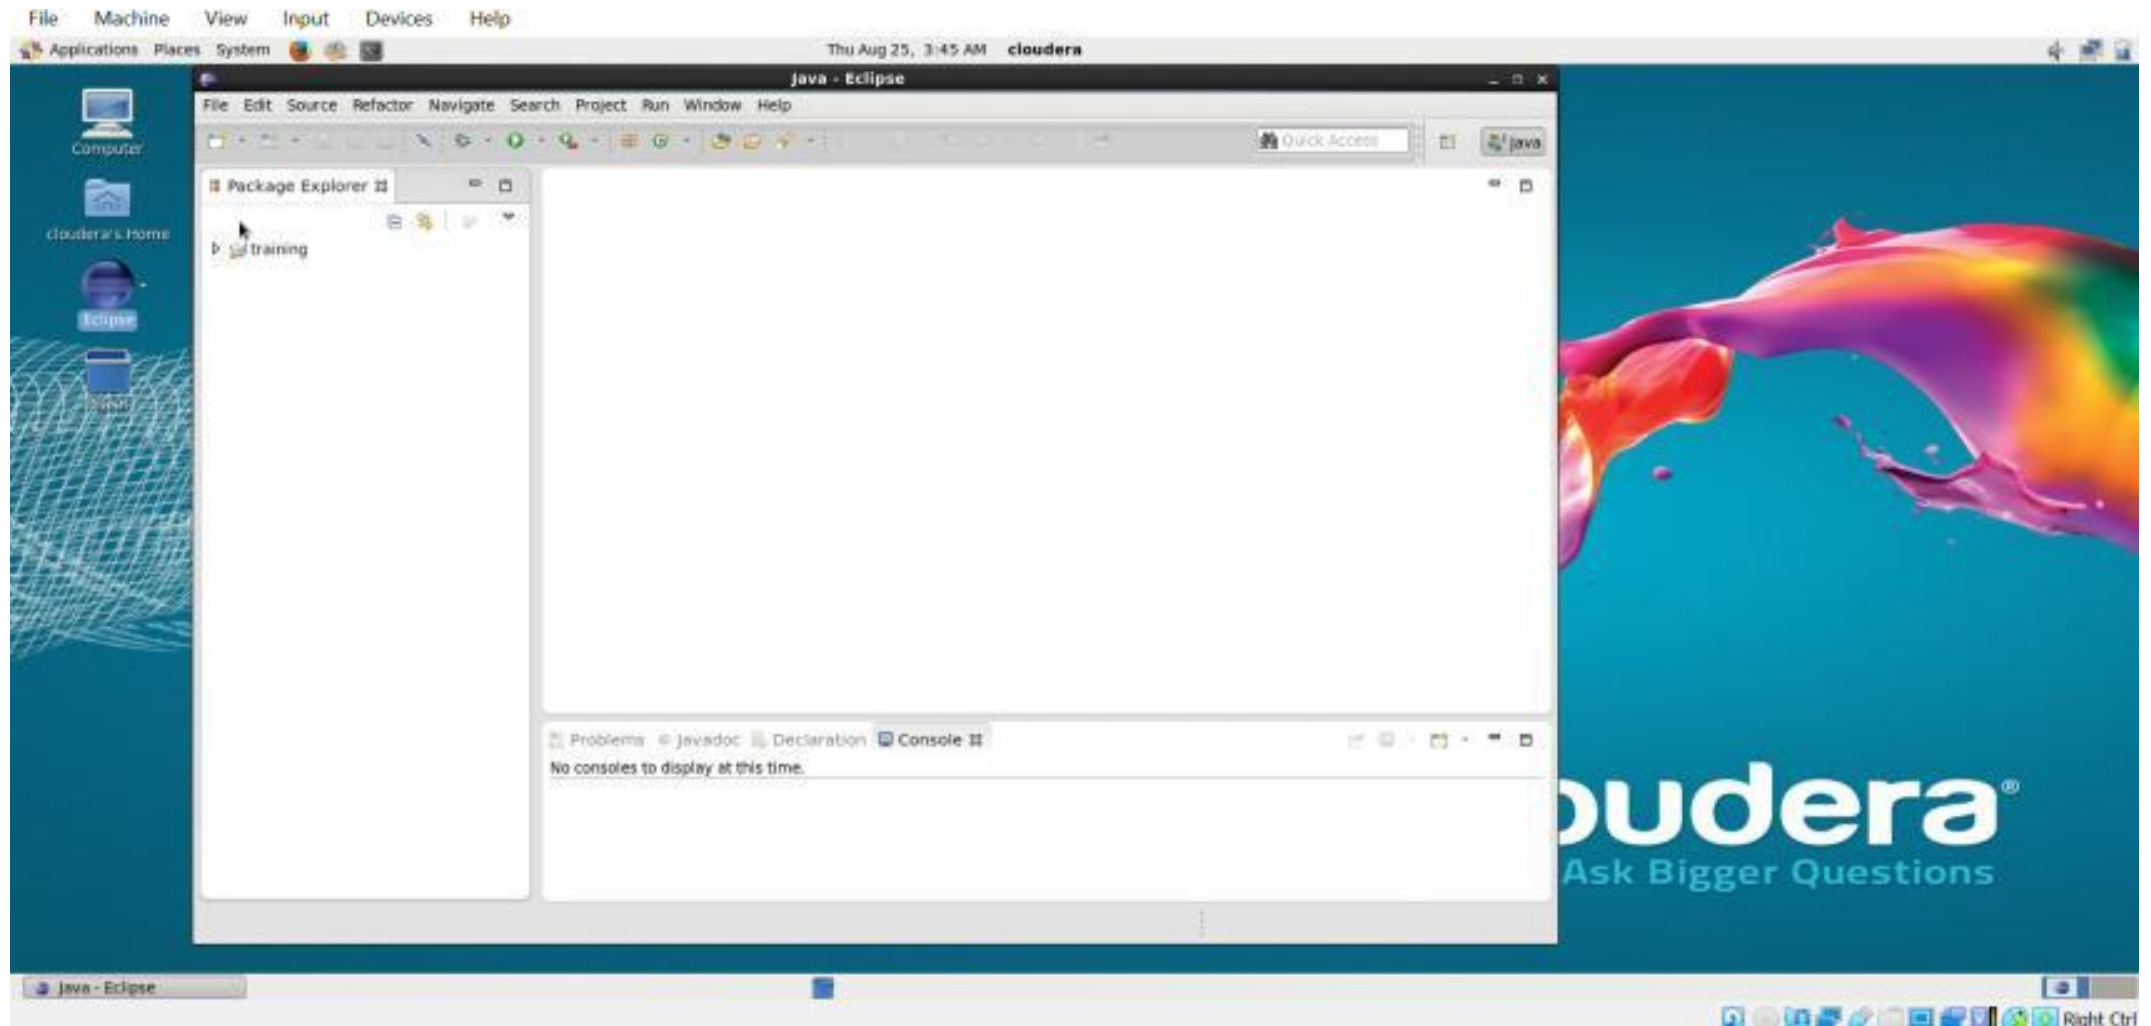


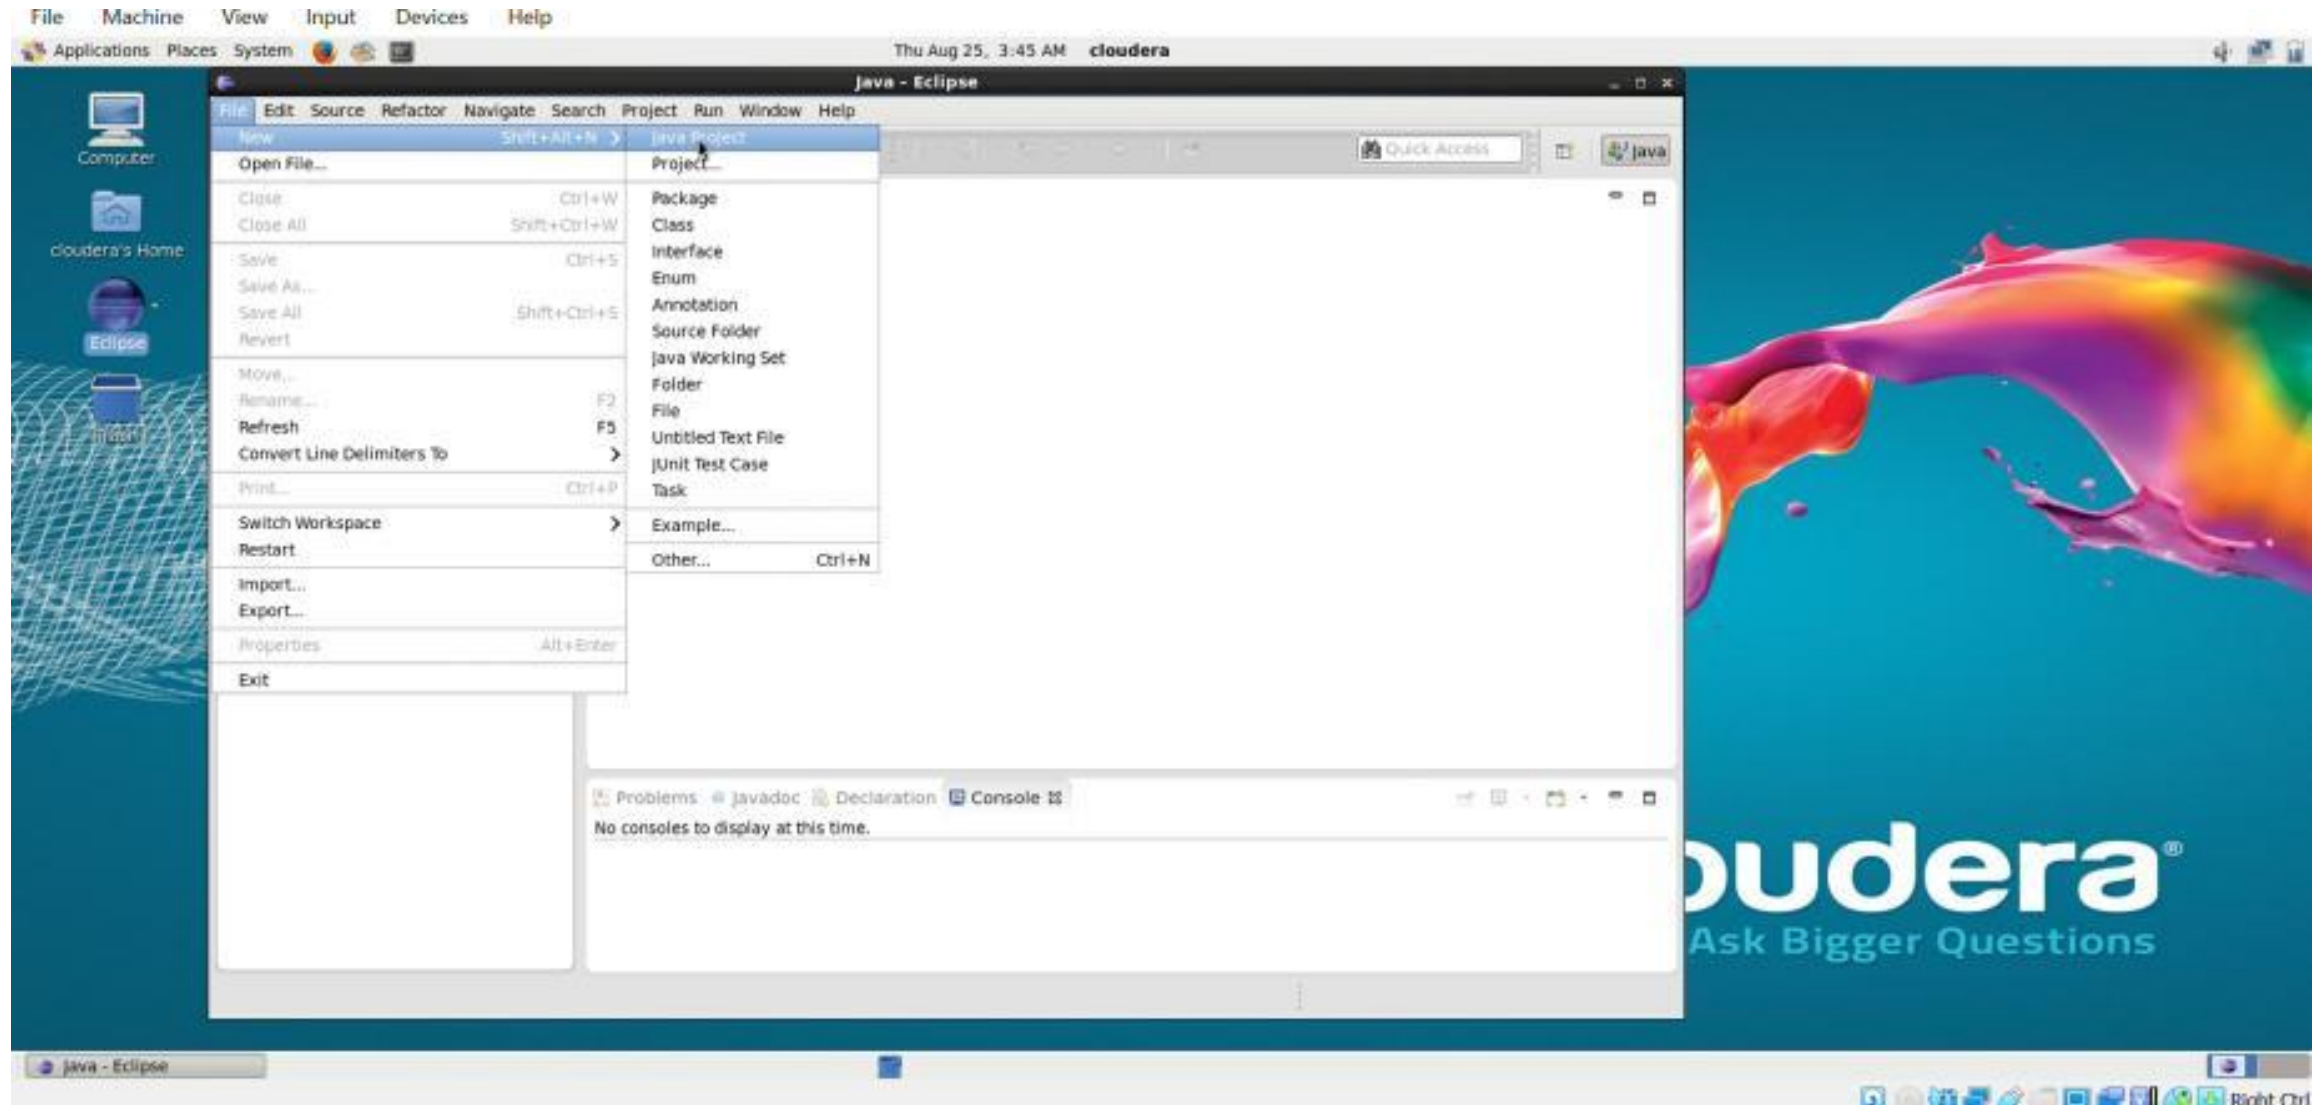
Aim:Run a basic Word Count MapReduce program to understand MapReduce Paradigm.

☐Open Oracle VM Virtual box -> click start -> open cloudera -> open eclipse.

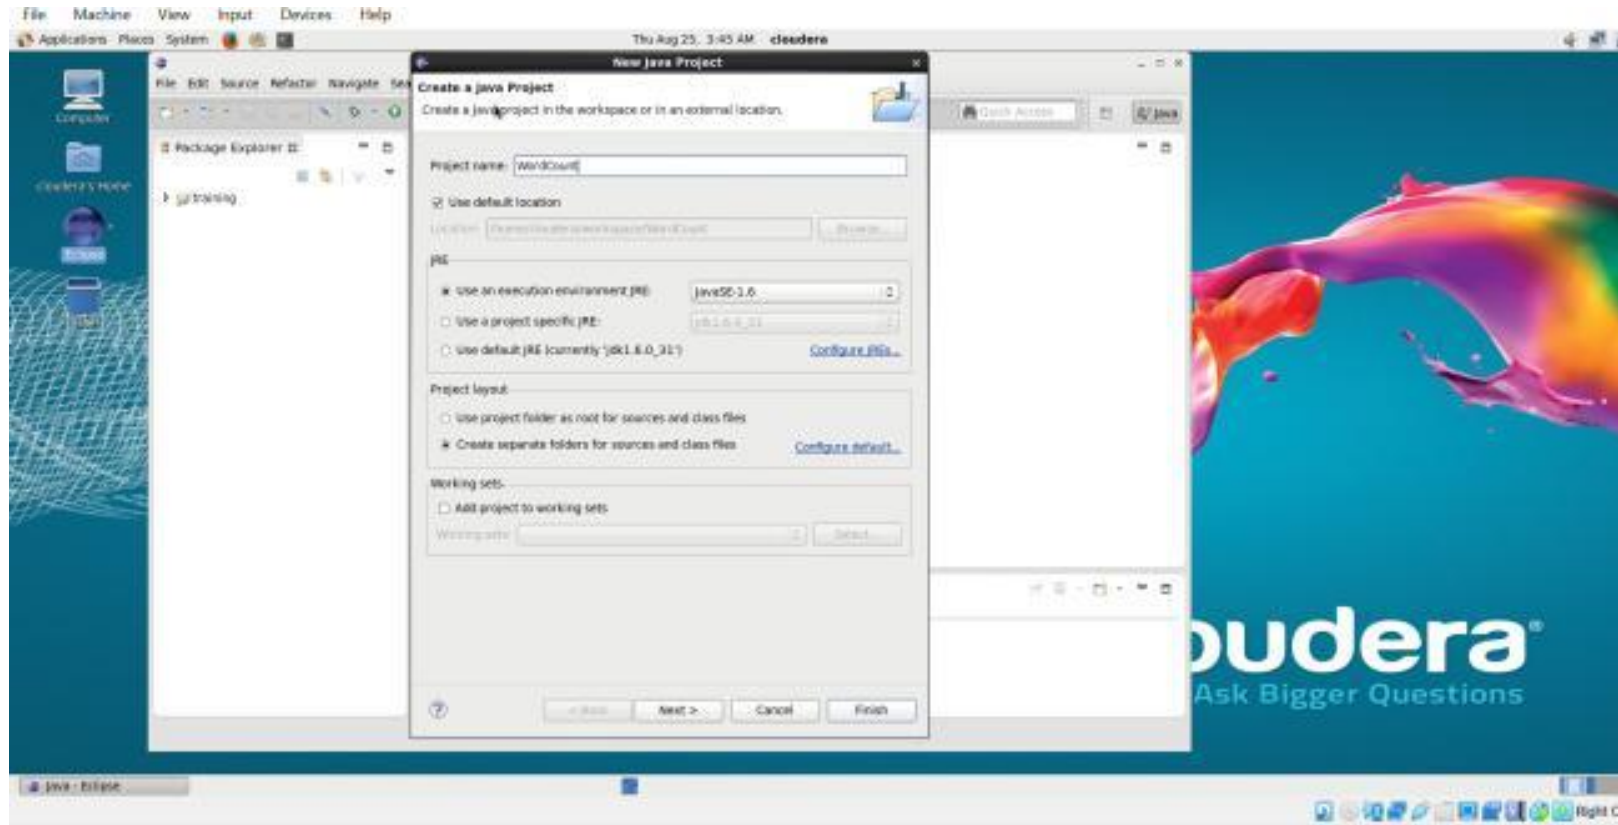




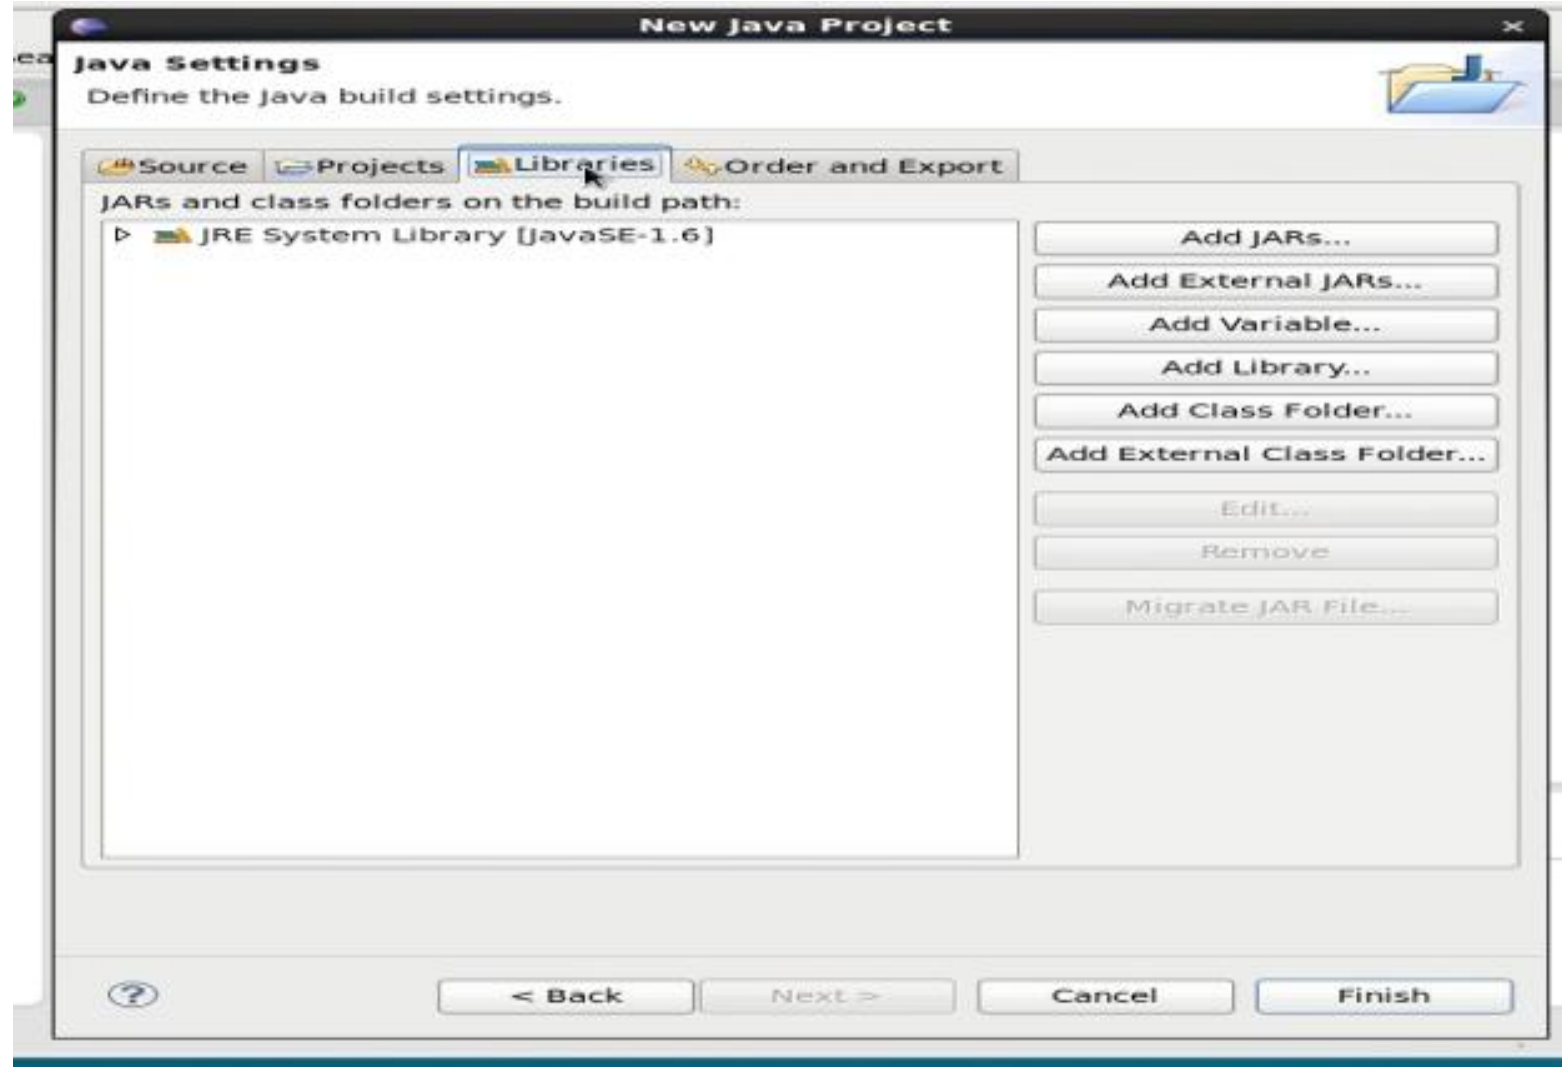
Click on file->New->Java Project



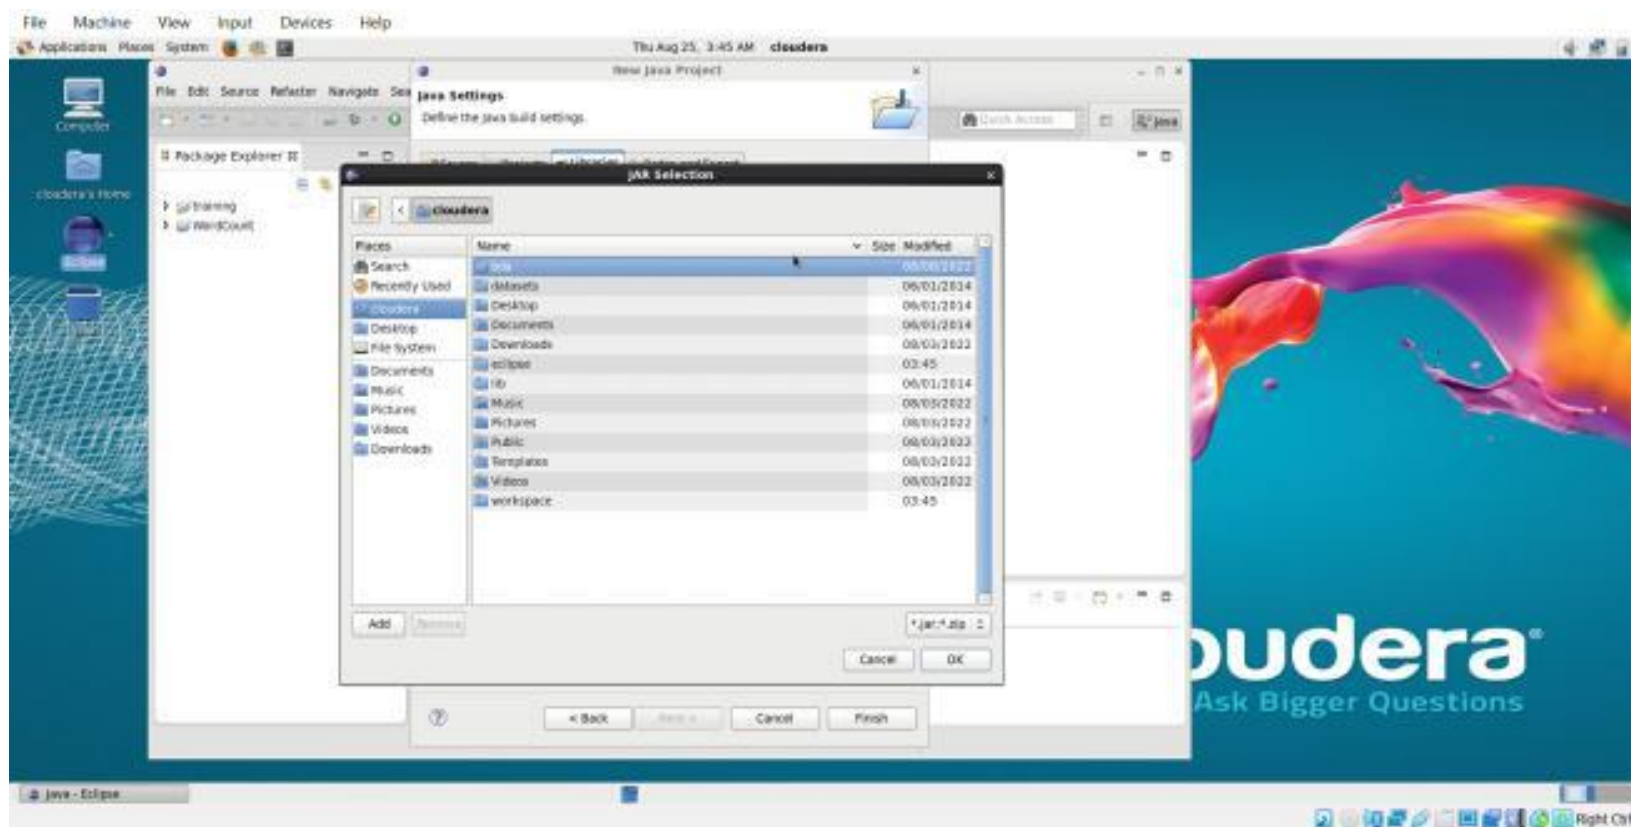
Give your project name as “WordCount” and click on next.



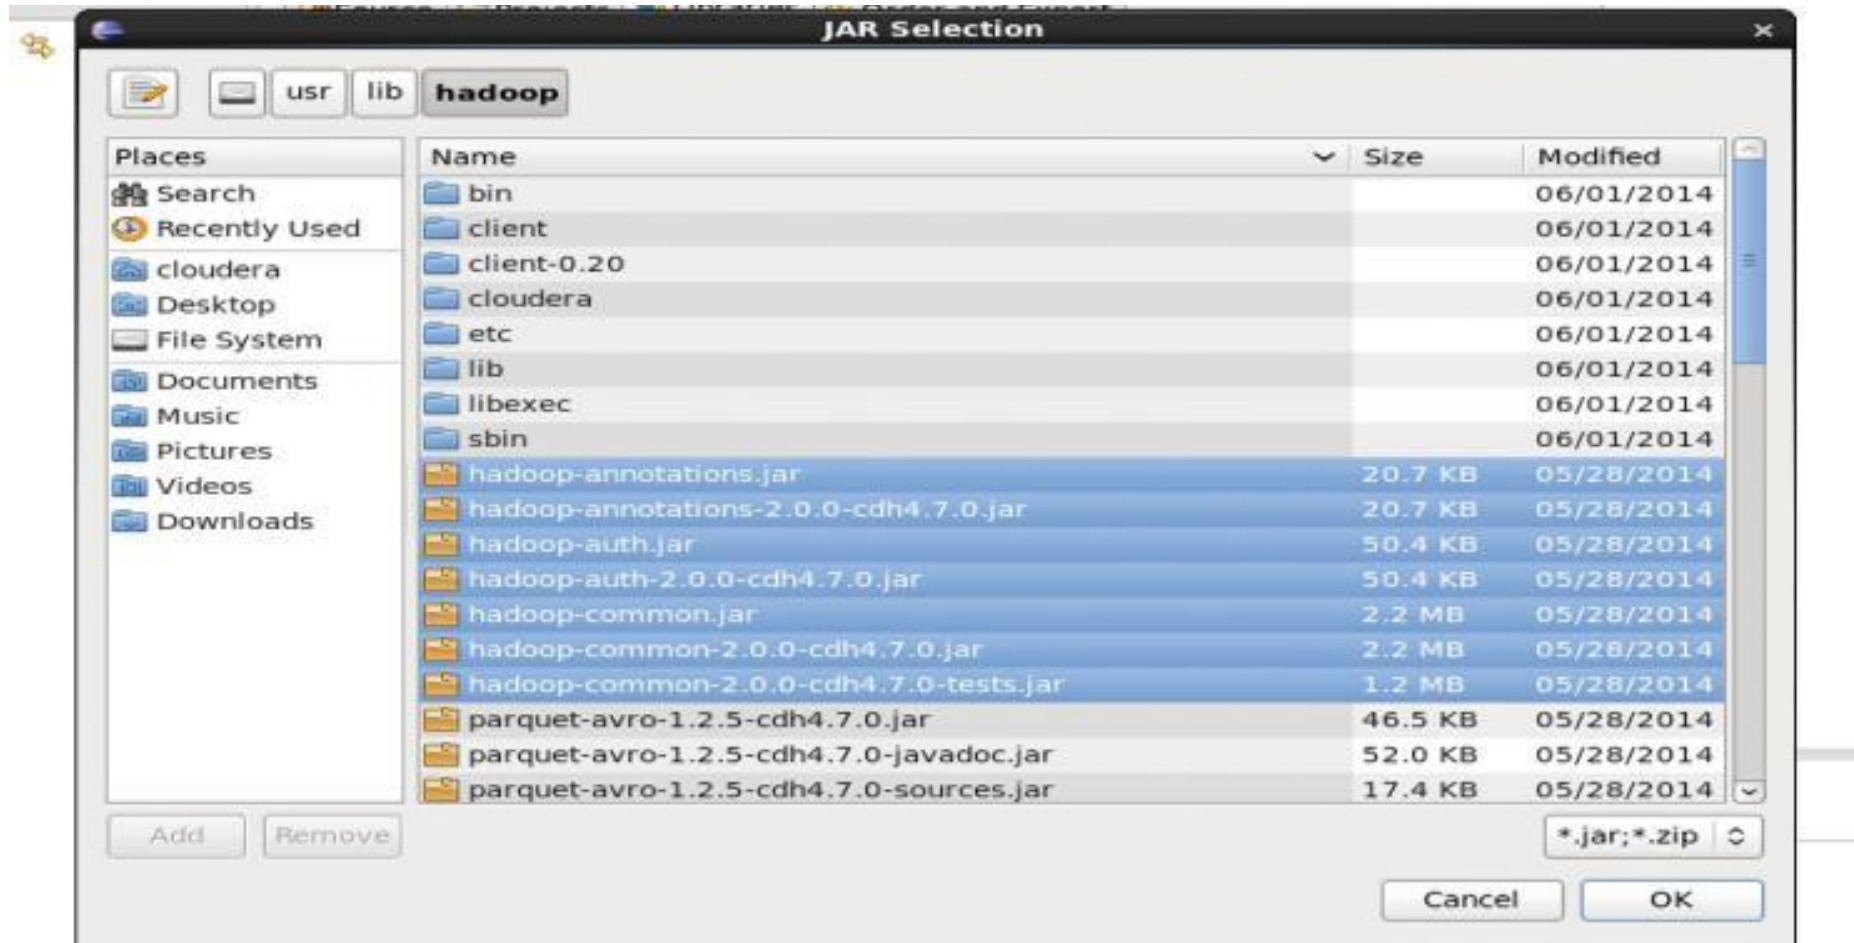
After that, click on libraries on the top and click on add external jars on the right side



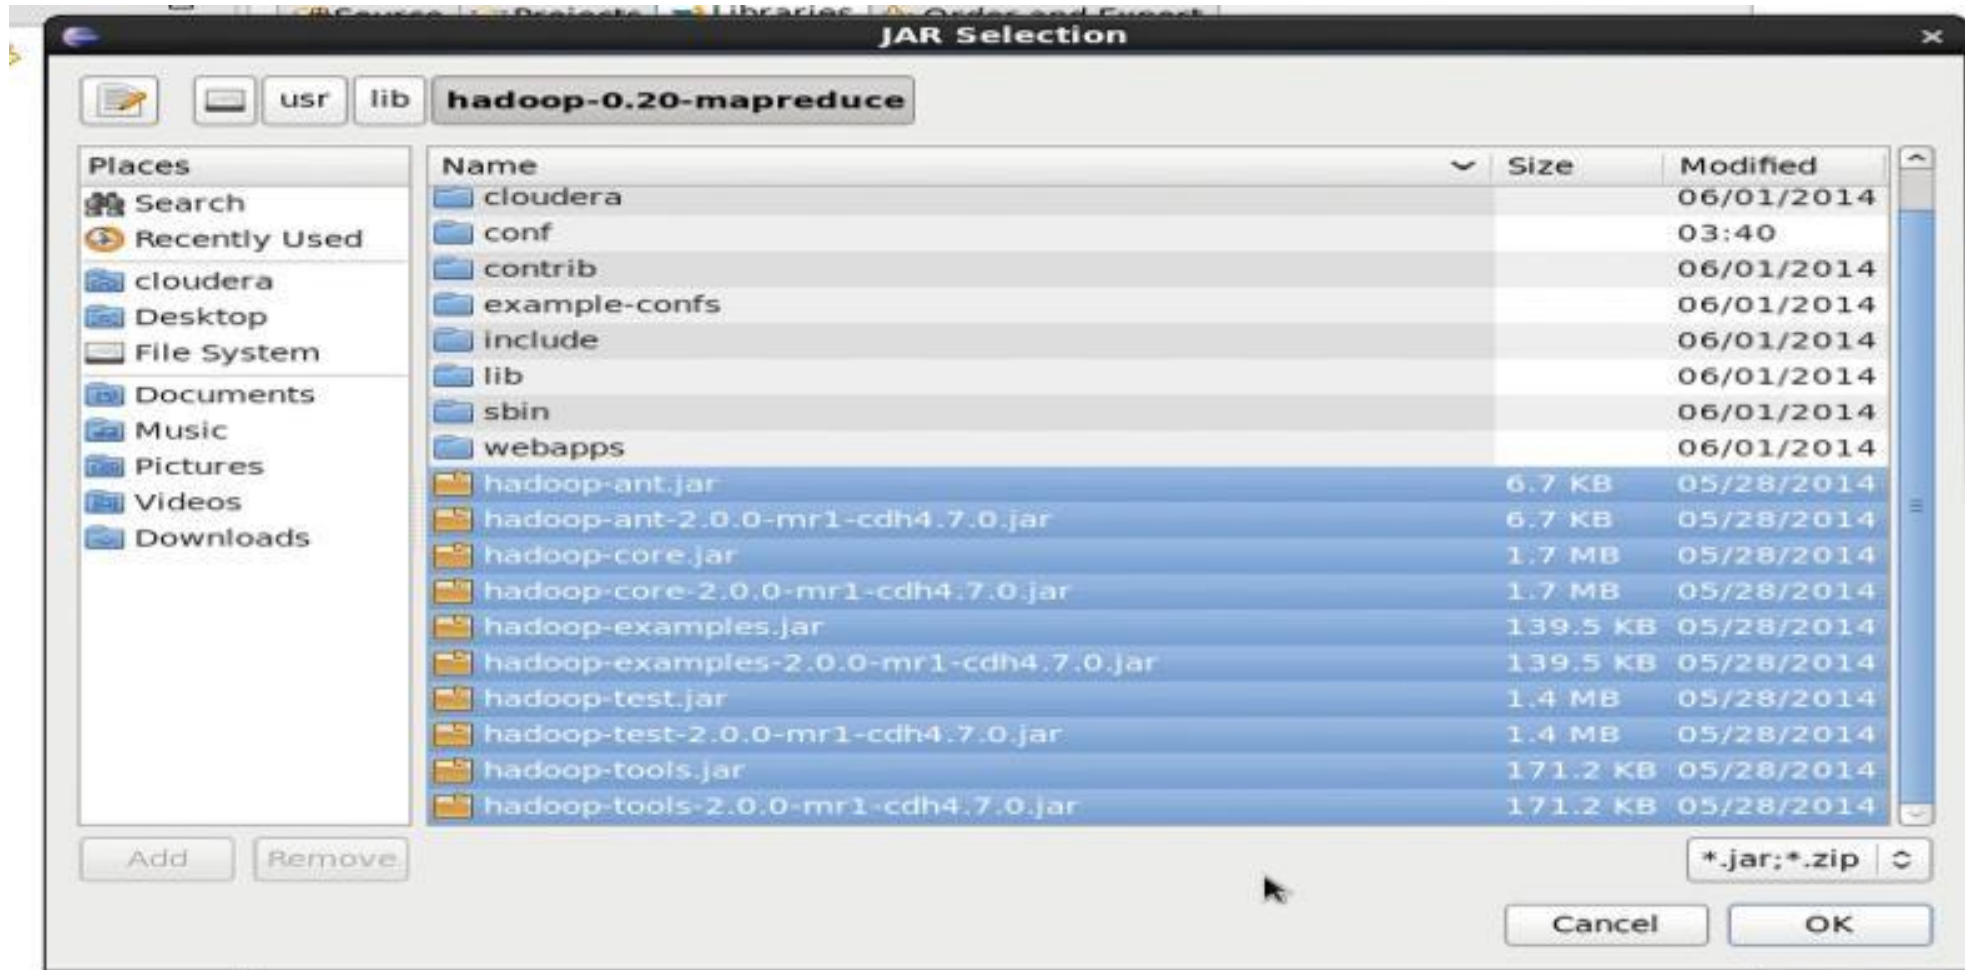
A dialog box will appear.




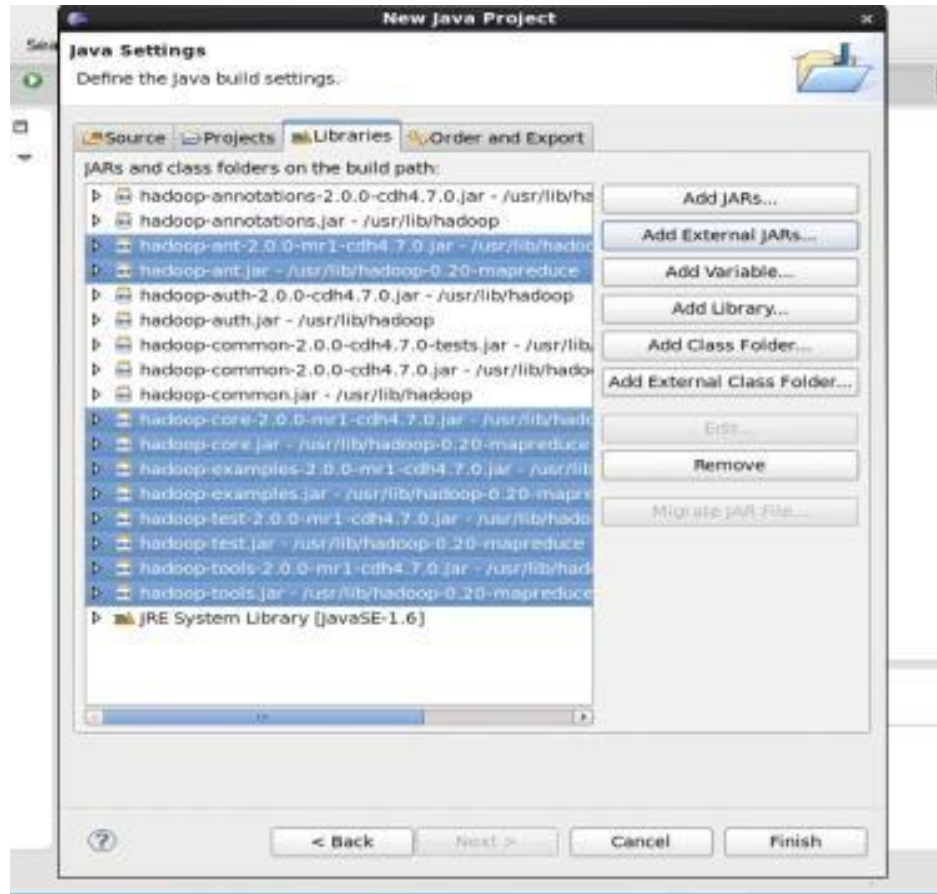
Click on file system on the left, and then usr->lib
->hadoop.
Select all the files named with hadoop and click ok.



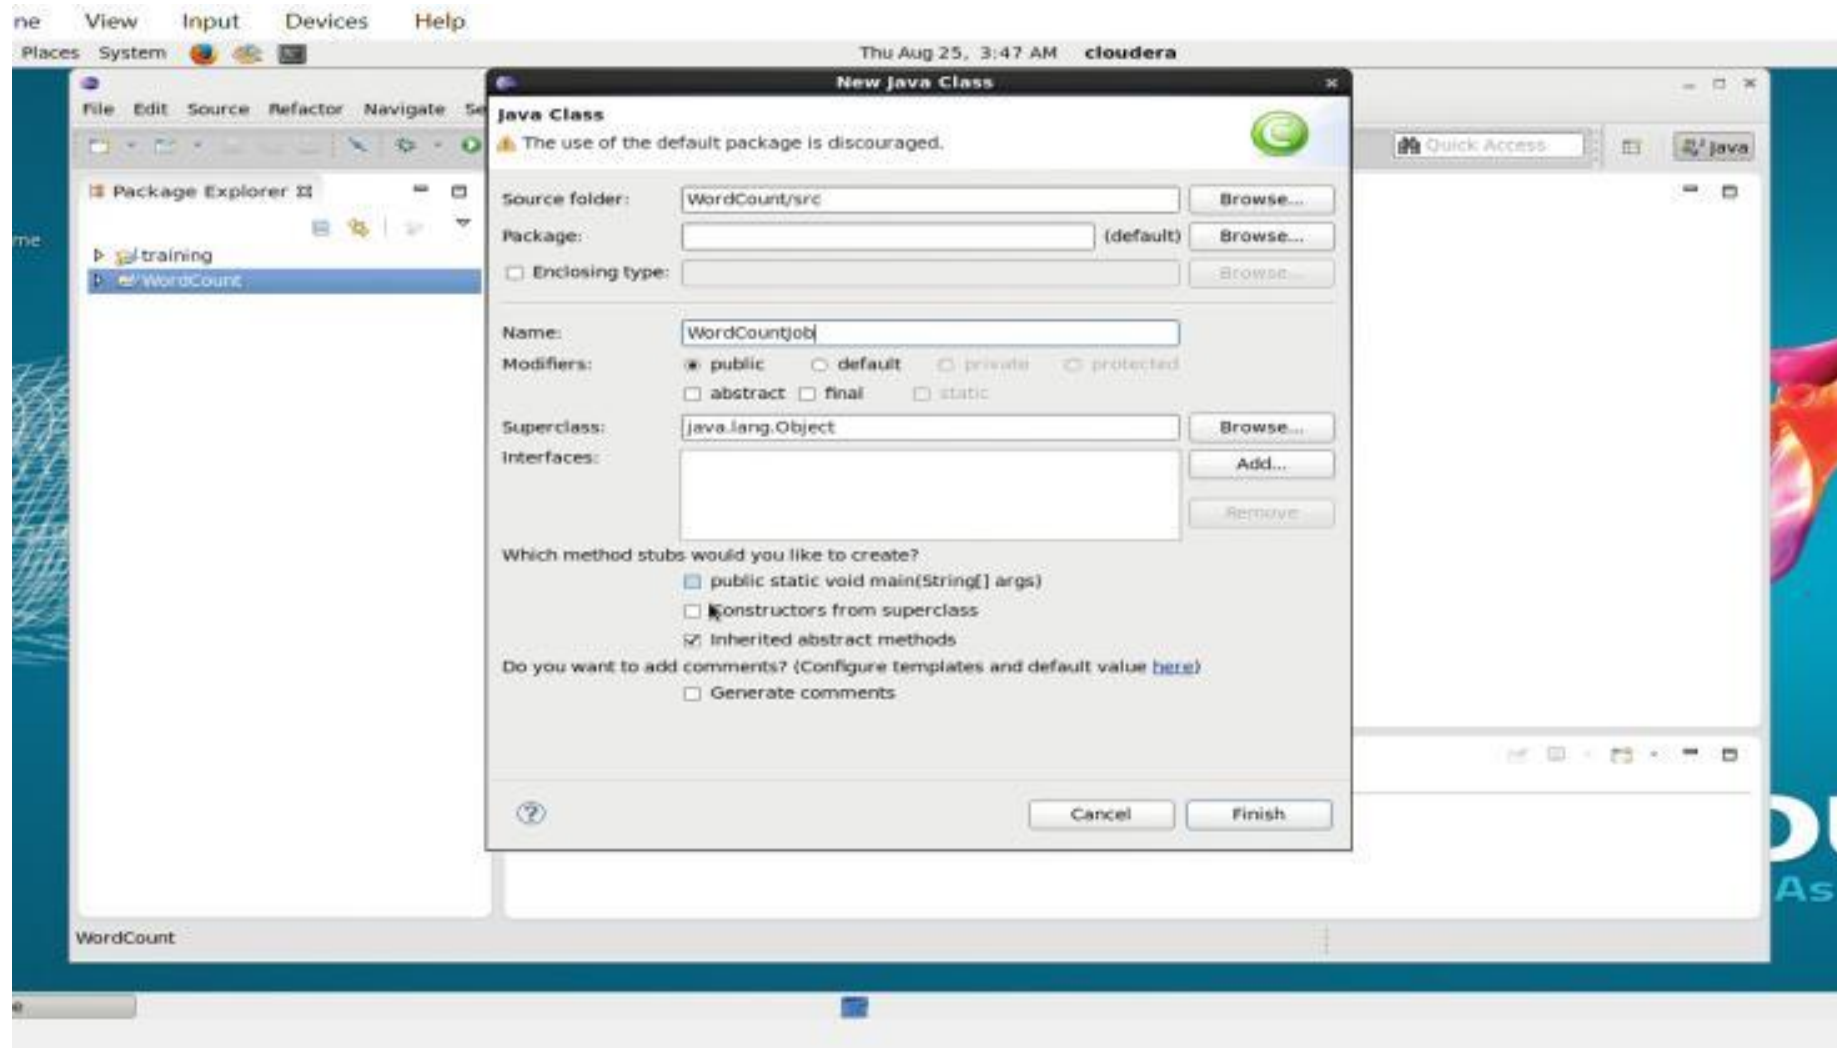
Go back to lib and click **hadoop-0.20-mapreduce** and select all the files named with hadoop and click ok.



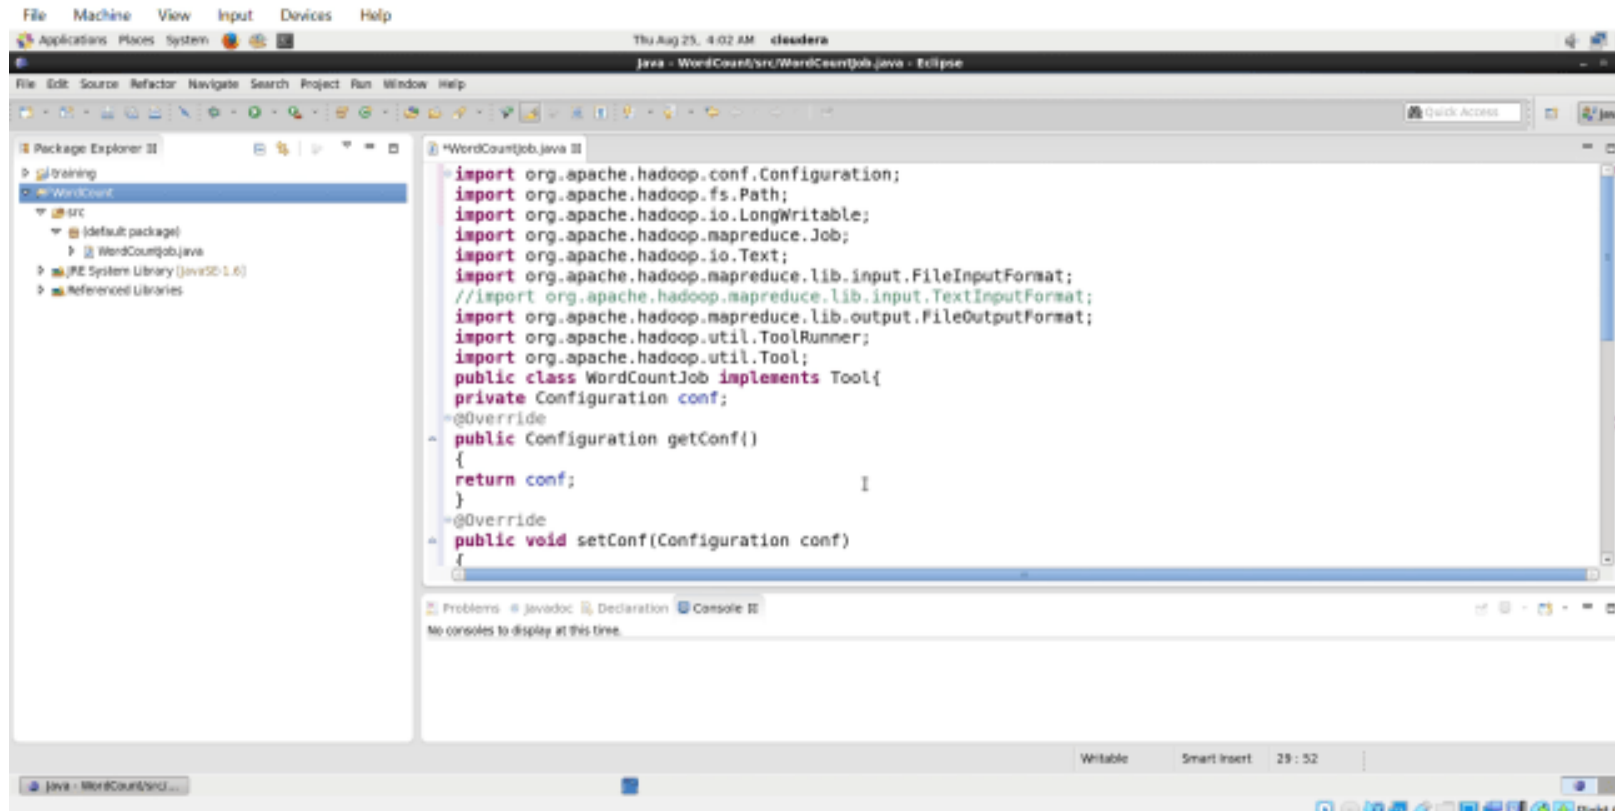
All the external files are added and now click on finish. The project is created successfully.



Create a class with name “WordCountJob”



Copy and paste the program 4 lines at a time-
If copy and paste doesn't work then goto device in
cloudera and set both shared clipboard and drag
and drop as bidirectional.



```
File  Machine  View  Input  Devices  Help
Thu Aug 25, 4:02 AM  cloudera
Java - WordCount/src/WordCountJob.java - Eclipse
File  Edit  Source  Refactor  Navigate  Search  Project  Run  Window  Help

Package Explorer
> training
  > WordCount
    > src
      > (default package)
        > WordCountJob.java
  > JRE System Library [jvse1.6]
  > Referenced Libraries

WordCountJob.java
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
//import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.ToolRunner;
import org.apache.hadoop.util.Tool;
public class WordCountJob implements Tool{
    private Configuration conf;

    @Override
    public Configuration getConf()
    {
        return conf;
    }

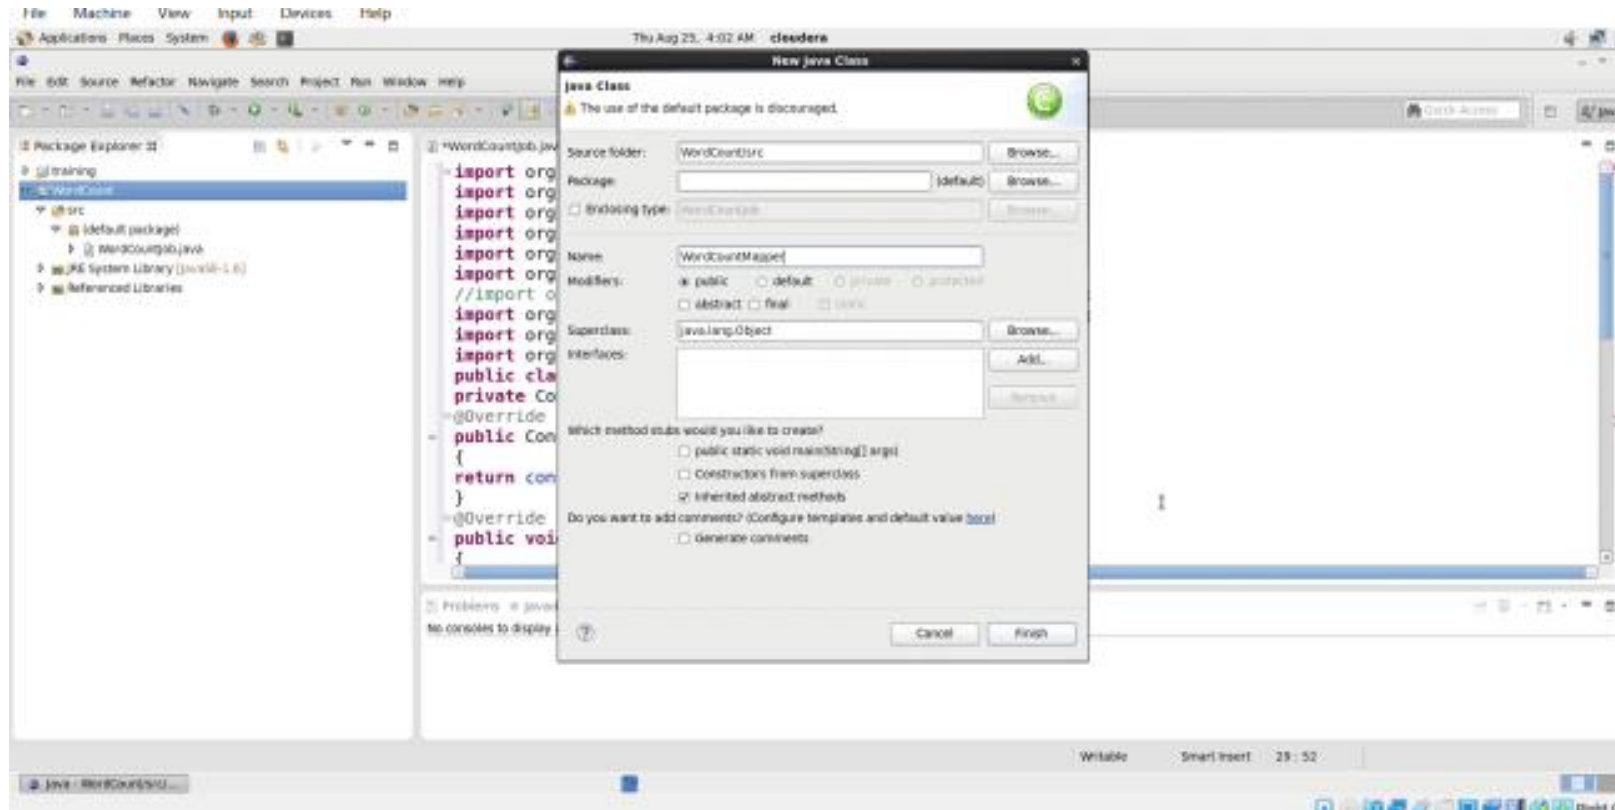
    @Override
    public void setConf(Configuration conf)
    {
        this.conf = conf;
    }
}
```

Problems Javadoc Declaration Console
No consoles to display at this time.

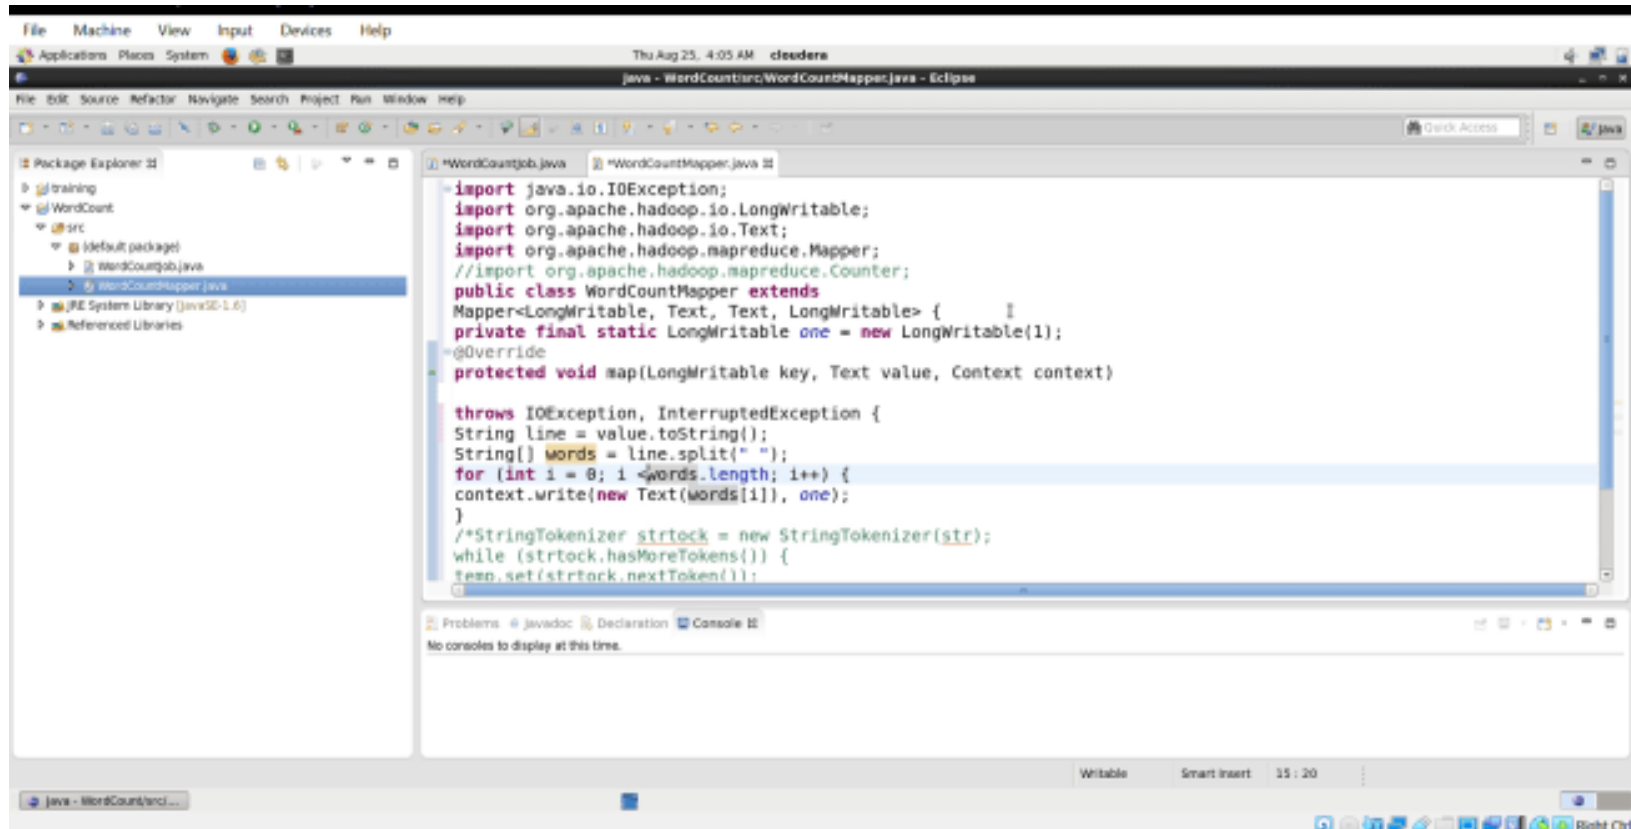
Java - WordCount/src/...

Writable Smart Insert 29 : 52

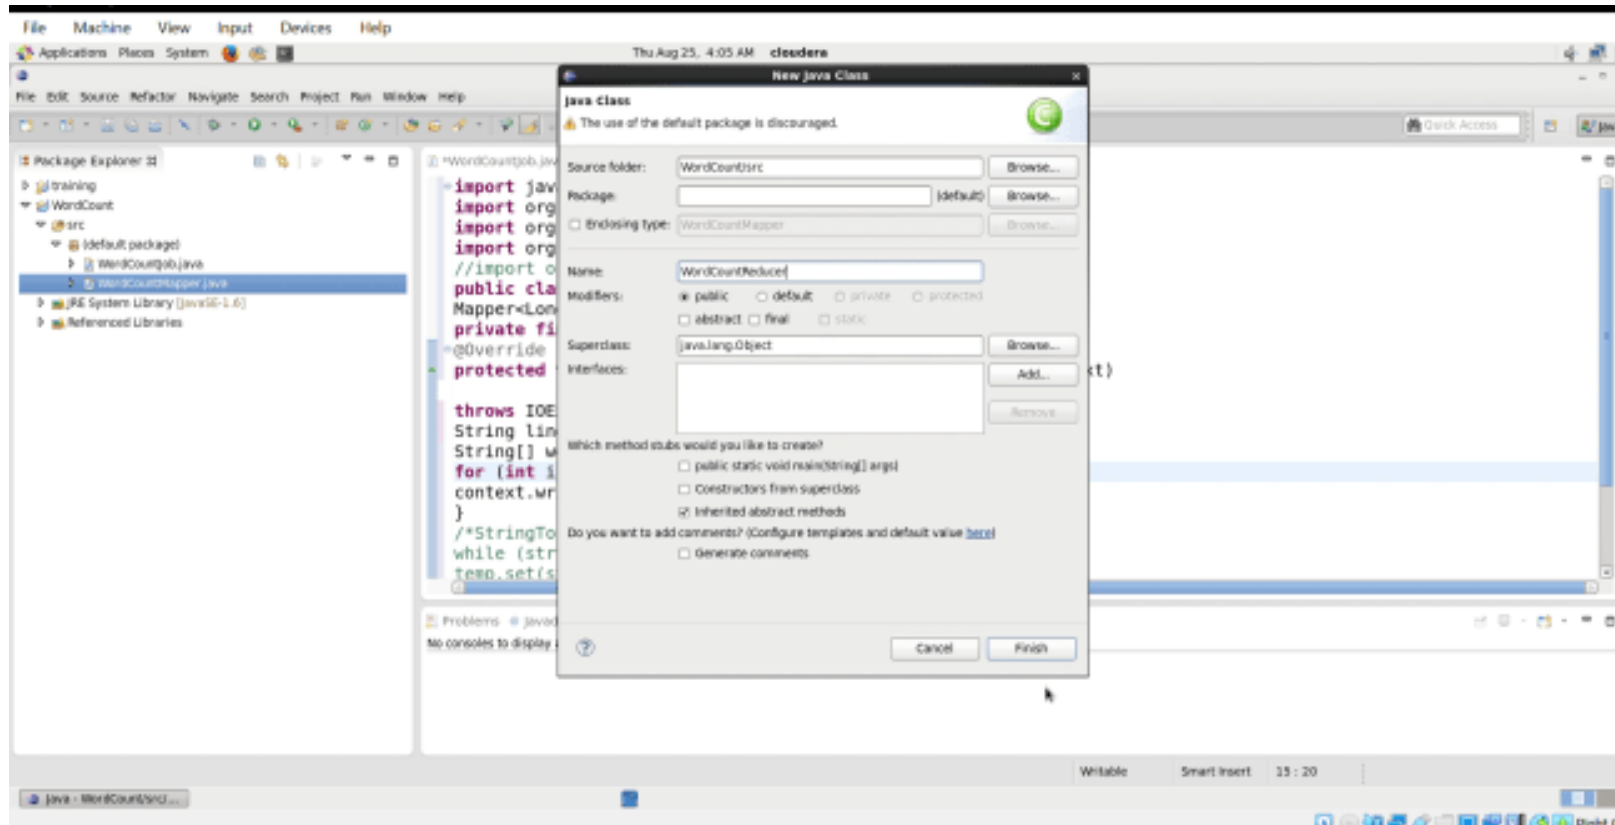
Create another class “WordCountMapper”



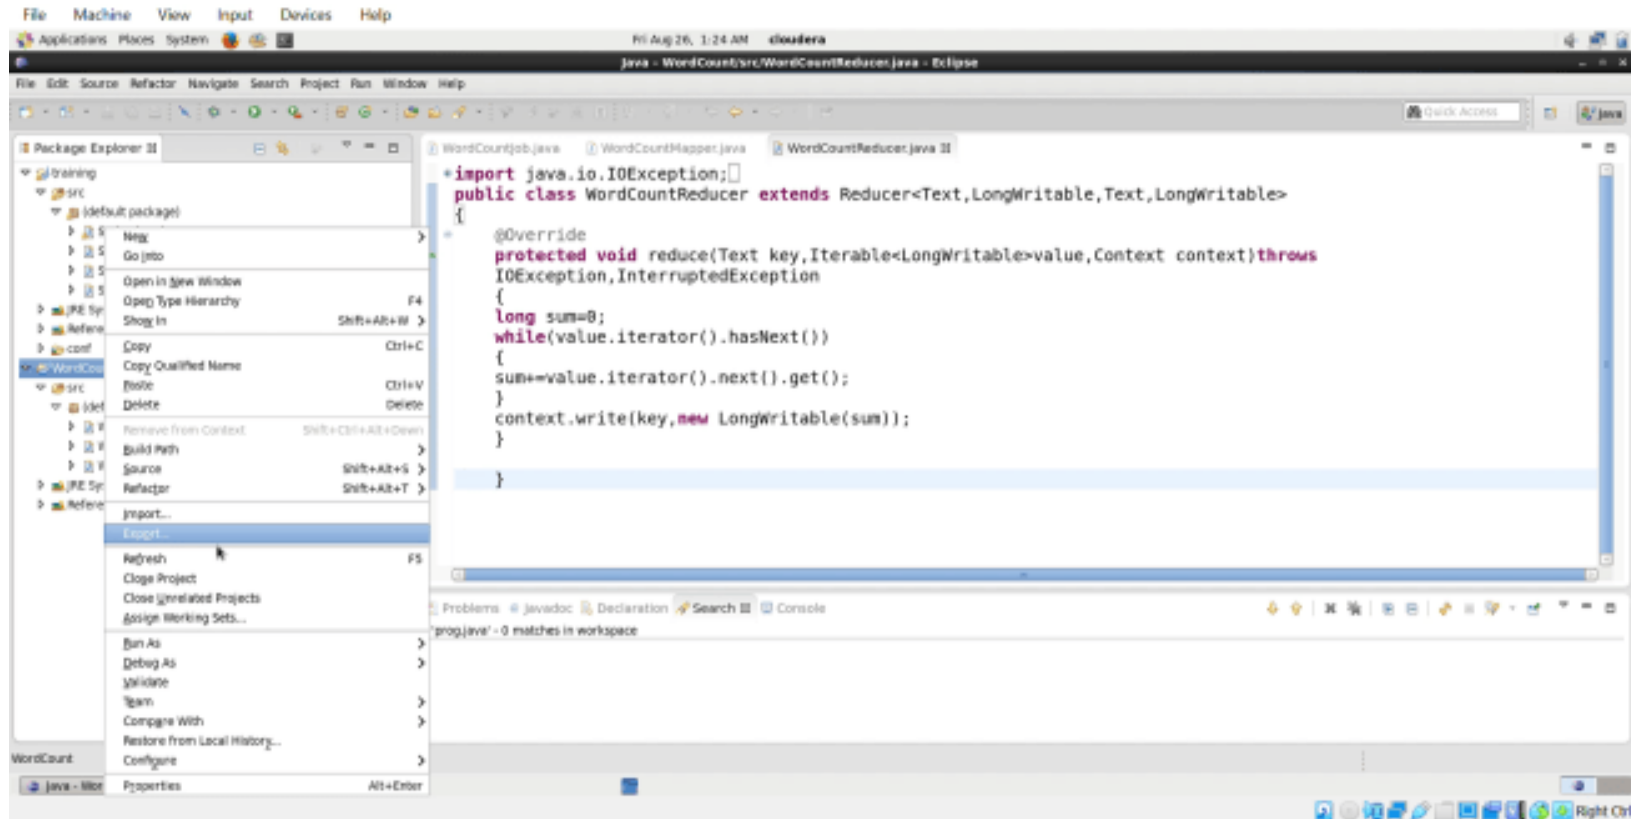
Copy and paste the program



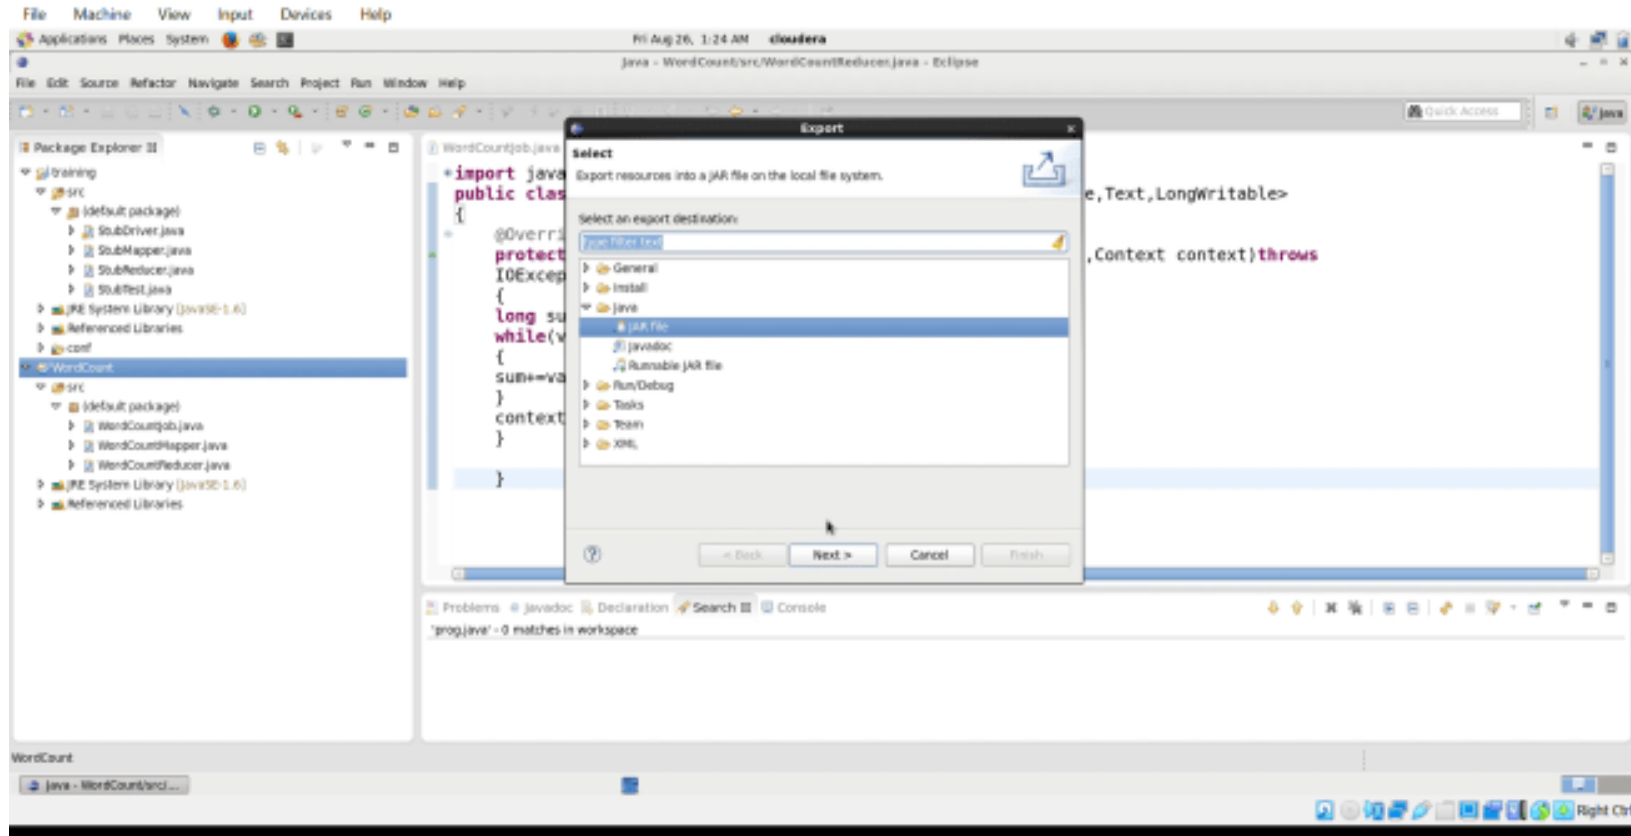
Create another class “WordCountReducer”



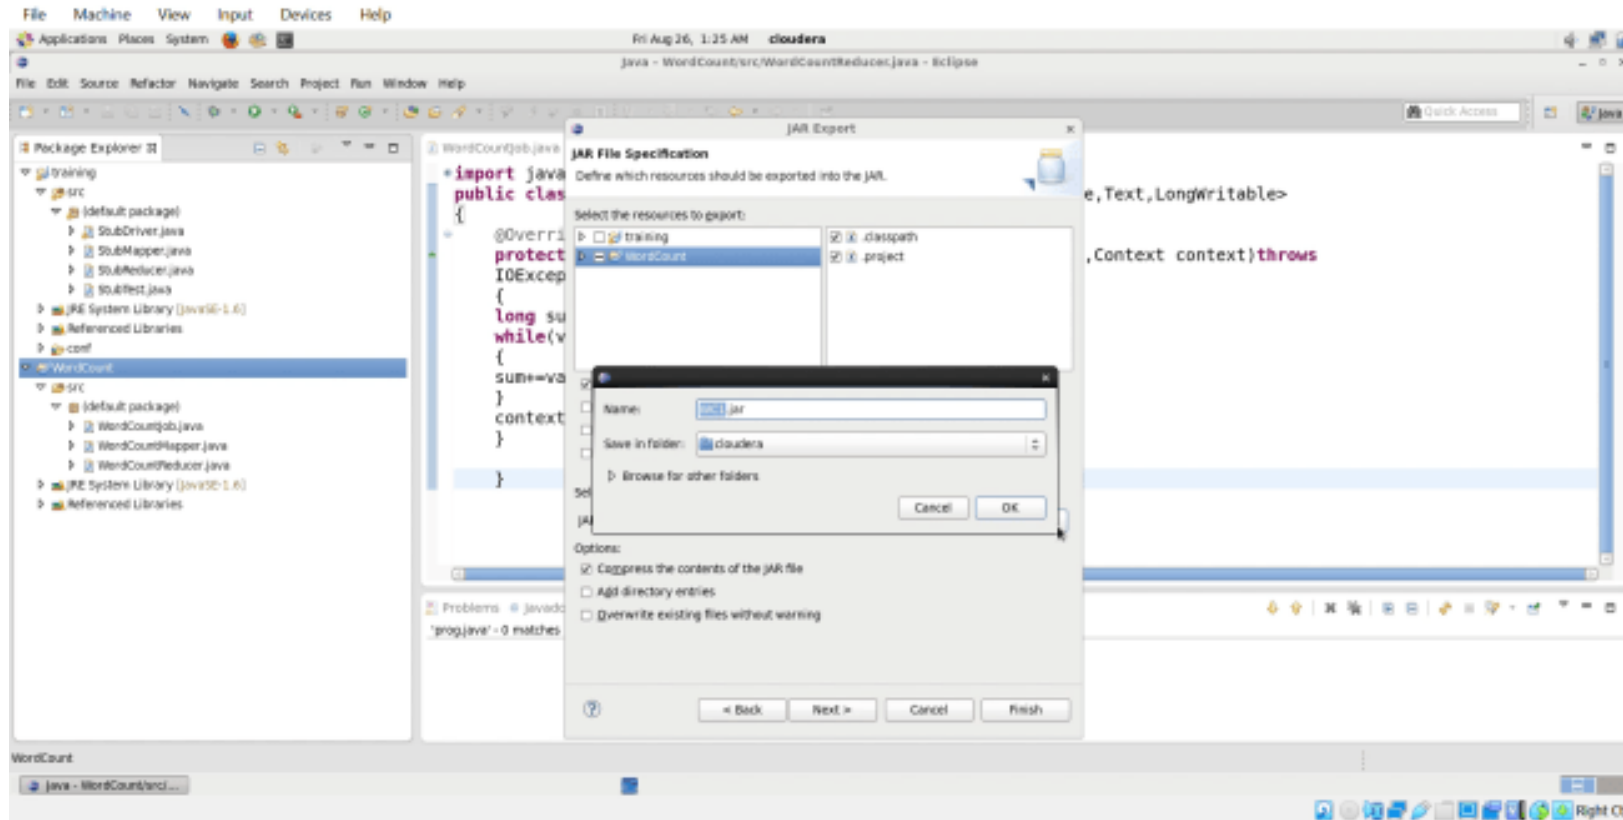
After finishing all the programs, right click on your project and select export.



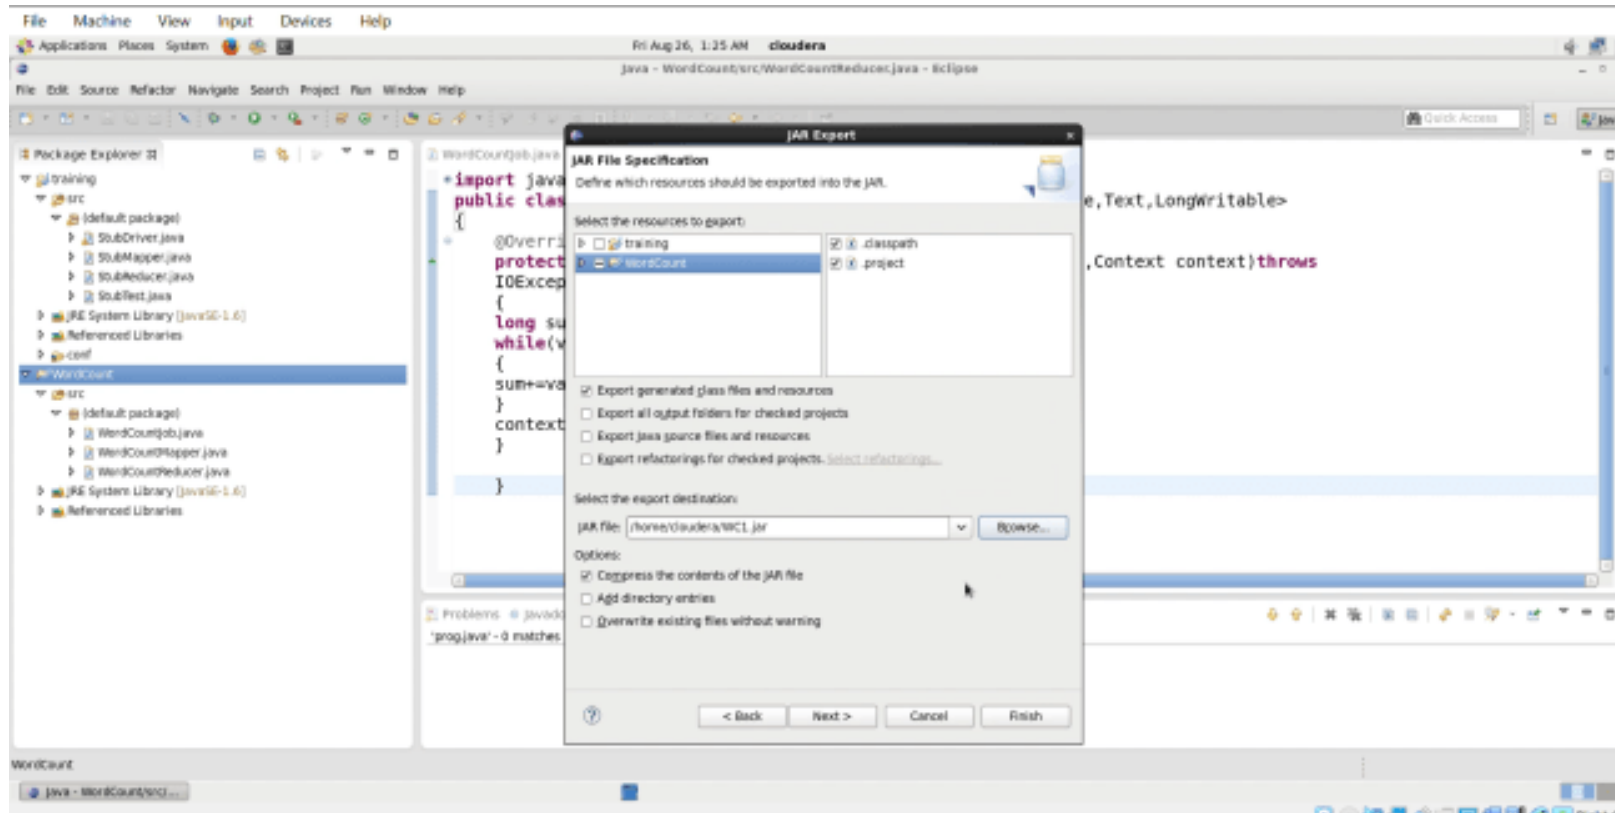
Select java-> JAR file->next



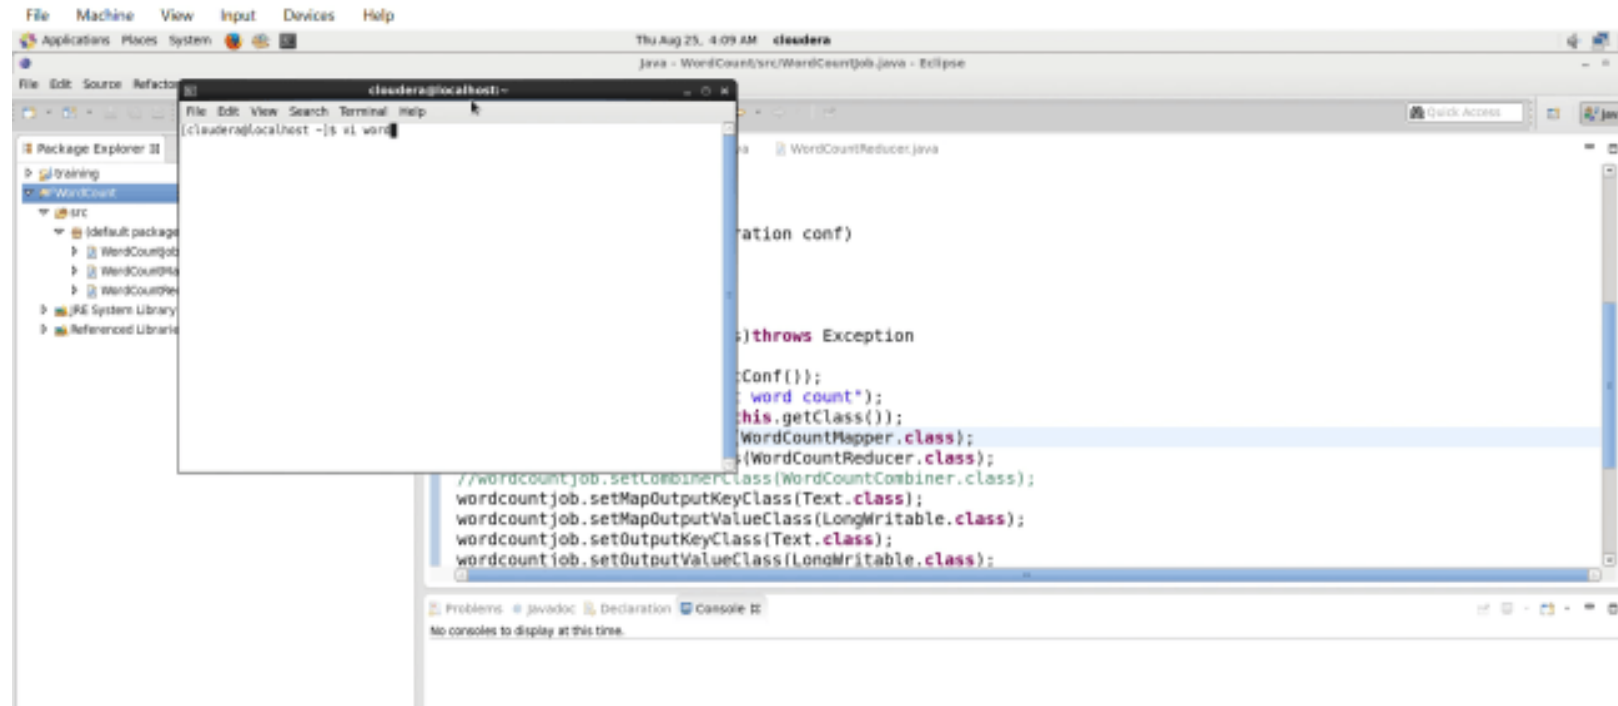
Click on browse and a dialog box will appear, give any name for the jar file and click ok.



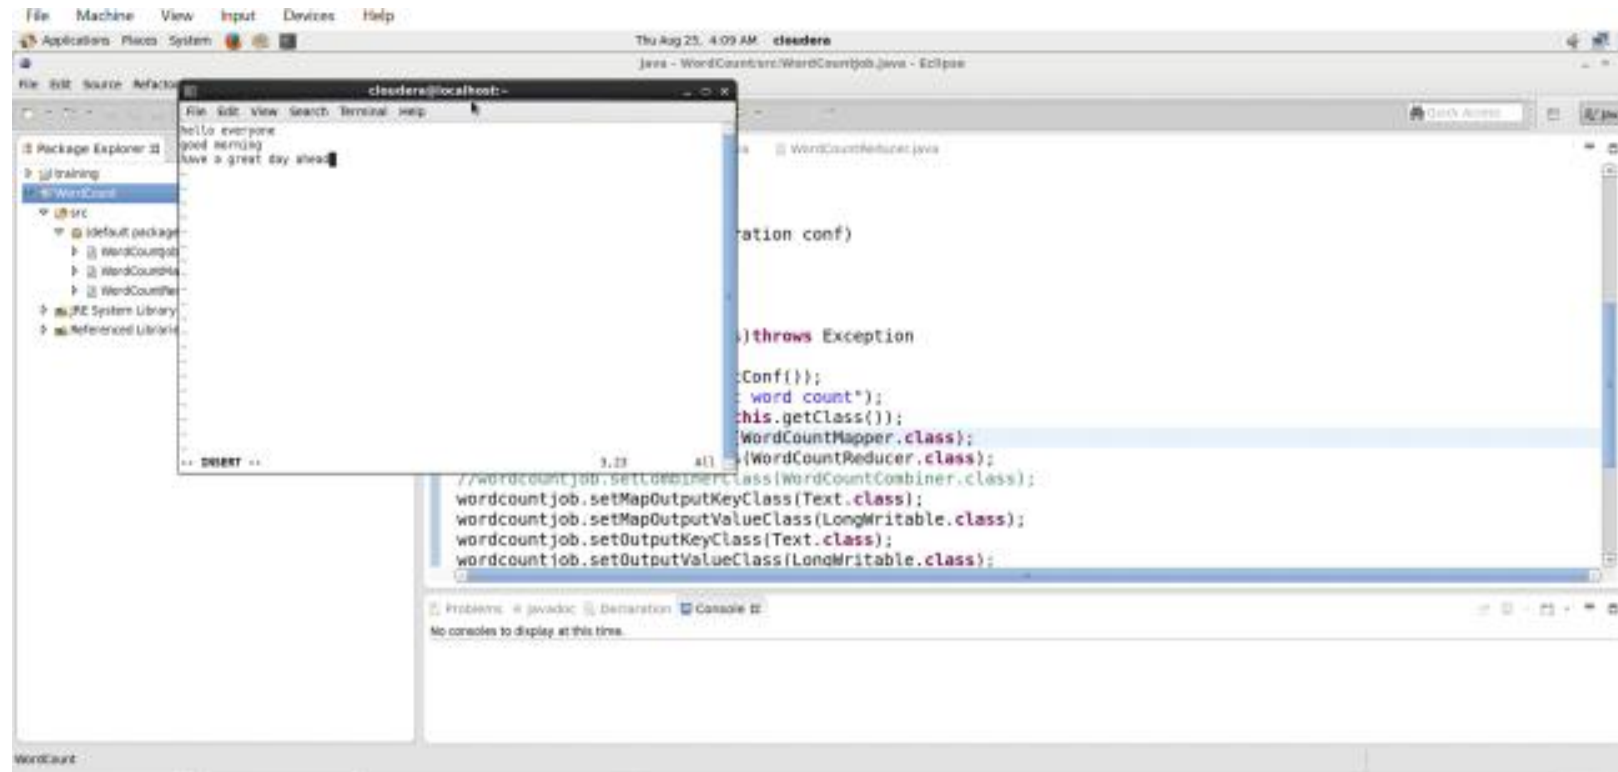
Click on finish.



Now, open the terminal and create a text file “word”.



Insert few sentences in the file.



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ ls  
bda Desktop eclipse Music Templates WC.jar  
count Documents file Pictures Videos word.txt  
datasets Downloads lib Public WC1.jar workspace  
[cloudera@localhost ~]$ hadoop fs -ls  
Found 3 items  
drwx----- - cloudera cloudera 0 2022-08-26 00:35 .Trash  
drwx----- - cloudera cloudera 0 2022-08-26 00:29 .staging  
-rw-r--r-- 3 cloudera cloudera 51 2022-08-25 04:21 word.txt  
[cloudera@localhost ~]$ hadoop fs -cat word.txt  
hello everyone  
good morning  
have a great day ahead  
[cloudera@localhost ~]$ hadoop jar WC1.jar WordCountJob word.txt result  
22/08/26 01:27:12 WARN mapred.JobClient: Use GenericOptionsParser for parsing the  
arguments. Applications should implement Tool for the same.  
22/08/26 01:27:12 INFO input.FileInputFormat: Total input paths to process : 1  
22/08/26 01:27:13 INFO mapred.JobClient: Running job: job_202208260059_0001  
22/08/26 01:27:14 INFO mapred.JobClient: map 0% reduce 0%  
22/08/26 01:27:27 INFO mapred.JobClient: map 100% reduce 0%
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
22/08/26 01:27:41 INFO mapred.JobClient: Physical memory (bytes) snapshot=385183744  
22/08/26 01:27:41 INFO mapred.JobClient: Virtual memory (bytes) snapshot=2013032448  
22/08/26 01:27:41 INFO mapred.JobClient: Total committed heap usage (bytes)=232067072  
[cloudera@localhost ~]$ hadoop fs -ls result  
Found 4 items  
-rw-r--r-- 3 cloudera cloudera 0 2022-08-26 01:27 result/_SUCCESS  
drwxr-xr-x - cloudera cloudera 0 2022-08-26 01:27 result/_logs  
-rw-r--r-- 3 cloudera cloudera 38 2022-08-26 01:27 result/part-r-00000  
-rw-r--r-- 3 cloudera cloudera 31 2022-08-26 01:27 result/part-r-00001  
[cloudera@localhost ~]$ hadoop fs -cat /user/cloudera/result/part-r-00000  
a 1  
ahead 1  
everyone 1  
good 1  
great 1  
[cloudera@localhost ~]$ hadoop fs -cat /user/cloudera/result/part-r-00001  
day 1  
have 1  
hello 1  
morning 1  
[cloudera@localhost ~]$
```

OUTPUT:

The screenshot shows a web browser window titled "HDFS - Mozilla Firefox" with the address bar displaying "localhost:localdomain:50075/browseDirectory.jsp?nameNodeInfoPort=50070&srcAddr=localhost:8028". The page content shows the "Contents of directory /" with a search bar and a table of files and directories.

Contents of directory /

Go to: go

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
hbase	dir				2014-06-01 16:16	rwxr-xr-x	hbase	hbase
mrk	dir				2014-06-01 16:16	rwxr-xr-x	mrk	mrk
tmp	dir				2014-06-01 16:16	rwxr-xr-x	hdfs	supergroup
test	dir				2014-06-01 16:17	rwxr-xr-x	hdfs	supergroup
src	dir				2014-06-01 16:15	rwxr-xr-x	hdfs	supergroup

[Go back to DFS home](#)

Local logs

[Log directory](#)

[Hadoop, 2012.](#)

The bottom of the screenshot shows a terminal window with the prompt "cloudera@localhost" and a taskbar with various application icons.

FileMachineViewInputDevicesHelp

ApplicationsPlacesSystemFri Aug 26, 1:29 AMcloudera

HDFS:/user/cloudera - Mozilla Firefox

FileEditViewHistoryBookmarksToolsHelp

HDFS:/user/clouderalocalhost.localdomain:50075/browseDirectory.jsp?dir=%2Fuser%2Fcloudera&go=go&namenodeInfoPort=50070&nnaddr=localhost%3A8020Google

Most VisitedClouderaCloudera ManagerHueHDFS NameNodeHadoop JobTrackerHBase MasterSolr

Contents of directory /user/cloudera

Goto : /user/cloudera go

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
.Trash	dir				2022-08-26 00:35	rwX----	cloudera	cloudera
.staging	dir				2022-08-26 01:27	rwX----	cloudera	cloudera
result	dir				2022-08-26 01:27	rwXr-xr-x	cloudera	cloudera
word.txt	file	51 B	3	128 MB	2022-08-25 04:21	rw-r--r--	cloudera	cloudera

[Go back to DFS home](#)

Local logs

[Log](#) directory

[Hadoop](#), 2022.

cloudera@focalhost: ~HDFS:/user/cloudera - ...

Quick Chat

File: [/user/cloudera/result/part-r-00000](#)Go to: [Go back to dir listing](#)[Advanced view/download options](#)

```
a      1
ahead  1
everyone      1
good      1
great      1
```

[Download this file](#)[Tail this file](#)Chunk size to view (in bytes, up to file's DFS block size):

Total number of blocks: 1

FileMachineViewInputDevicesHelp

ApplicationsPlacesSystemFri Aug 26, 1:29 AMcloudera

HDFS:/user/cloudera/result - Mozilla Firefox

FileEditViewHistoryBookmarksToolsHelp

HDFS:/user/cloudera/result

localhost.localdomain:50075/browseDirectory.jsp?dir=%2Fuser%2Fcloudera%2Fresult&namenodeinfoPort=50070&nnaddr=localhost:8020Google

Most VisitedClouderaCloudera ManagerHueHDFS NameNodeHadoop JobTrackerHBase MasterSolr

Contents of directory /user/cloudera/result

Goto : /user/cloudera/result go

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
SUCCESS	file	0 B	3	128 MB	2022-08-26 01:27	rw-r--r--	cloudera	cloudera
_logs	dir				2022-08-26 01:27	rw-r--r--	cloudera	cloudera
part-r-00000	file	38 B	3	128 MB	2022-08-26 01:27	rw-r--r--	cloudera	cloudera
part-r-00001	file	31 B	3	128 MB	2022-08-26 01:27	rw-r--r--	cloudera	cloudera

[Go back to DFS home](#)

Local logs

[Log](#) directory

[Hadoop](#), 2022.

localhost.localdomain:50075/browseDirectory.jsp?dir=/user/cloudera/result/part-r-00001&namenodeinfoPort=50070&nnaddr=localhost:8020

cloudera@localhost: ~HDFS:/user/cloudera/r...

Right Ctrl

FileMachineViewInputDevicesHelp

Fri Aug 26, 1:29 AMcloudera

ApplicationsPlacesSystem

HDFS:/user/cloudera/result/part-r-00001 - Mozilla Firefox

FileEditViewHistoryBookmarksToolsHelp

HDFS:/user/cloudera/result/part-r-00001

localhost.localdomain:50075/browseBlock.jsp?blockId=8484979633973710307&blockSize=31&genstamp=13436&filename=%2Fuser%2Fcloudera%2Fresult%2Fpart-r-00001&datanodePort=50010&nar

Google

Most Visited

ClouderaCloudera ManagerHueHDFS NameNodeHadoop JobTrackerHBase MasterSolr

File: /user/cloudera/result/part-r-00001

Goto: /user/cloudera/result go

[Go back to dir listing](#)
[Advanced view/download options](#)

day 1
have 1
hello 1
morning 1

[Download this file](#)
[Tail this file](#)

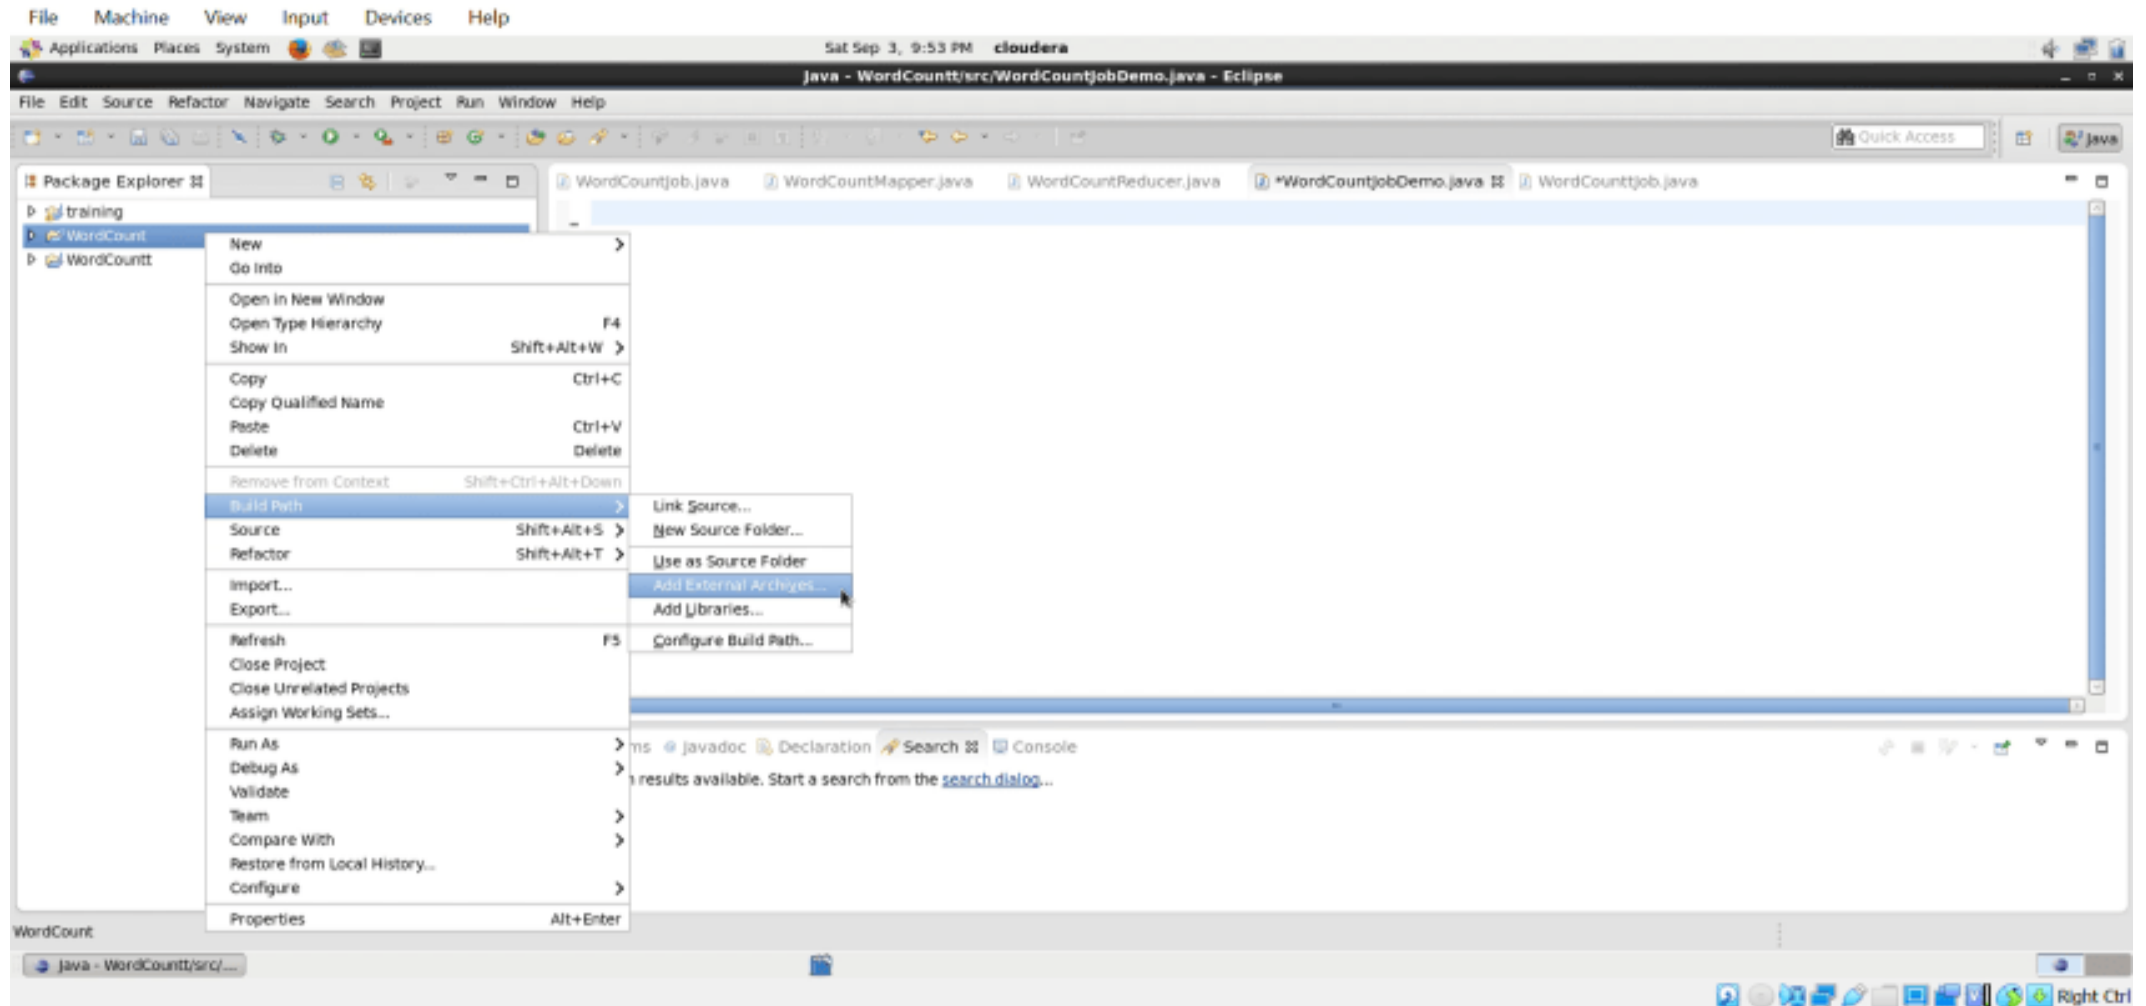
Chunk size to view (in bytes, up to file's DFS block size): 32768 Refresh

Total number of blocks: 1

cloudera@localhost: ~HDFS:/user/cloudera/r...

Right Ctrl

If there is any error in executing then set the path-
right click on your project->build path->add
external archives



Add all the hadoop and mapreduce external jar files and click ok.

