

BIG DATA ANALYTICS LAB

EXPERIMENT NO-05

AIM: Implement Matrix Multiplication with Hadoop Map Reduce.

Description:

- 1) Open Oracle VM VirtualBox->export cloudera->start
- 2) In eclipse->File->New->Java project->Project name “MatrixMultiplication”
->Finish

- 3) Create three classes.

Right click on MatrixMultiplication-> New->class->Name “MatrixDriver”

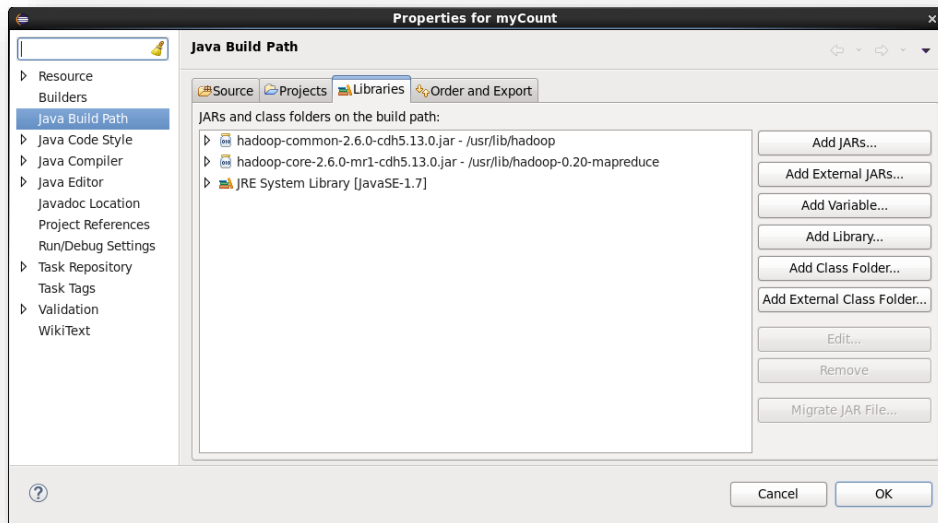
Right click on MatrixMultiplication-> New->class->Name “MatrixMapper”

Right click on MatrixMultiplication -> New->class->Name “MatrixReducer”

- 4) Add Hadoop libraries.

Right click on MatrixMultiplication->Build path->Configure Build path->Add external JARS. (usr\lib\hadoop\hadoop-common-2.6.0-cdh 5.13.0 jar,

usr\lib\hadoop\hadoop-core-2.6.0-cdh 5.13.0 jar)



PROGRAM:

MatrixDriver.java

```
import org.apache.hadoop.fs.Path;

import org.apache.hadoop.conf.*;

import org.apache.hadoop.io.*;

import org.apache.hadoop.mapreduce.*;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class MatrixDriver

{

    public static void main(String[] args) throws Exception

    {
```

```
Configuration conf = new Configuration();

// M is an m-by-n matrix; N is an n-by-p matrix.

conf.set("m", "2");

conf.set("n", "2");

conf.set("p", "2");

Job job = Job.getInstance(conf, "MatrixMultiplication");

FileInputFormat.addInputPath(job, new Path(args[0]));

FileOutputFormat.setOutputPath(job, new Path(args[1]));

job.setJarByClass(MatrixDriver.class);

job.setOutputKeyClass(Text.class);

job.setOutputValueClass(Text.class);

job.setMapperClass(MatrixMapper.class);

job.setReducerClass(MatrixReducer.class);

job.setInputFormatClass(TextInputFormat.class);

job.setOutputFormatClass(TextOutputFormat.class);

System.exit(job.waitForCompletion(true)?0:1);

}

}
```

MatrixMapper.java

```
import java.io.IOException;

import org.apache.hadoop.conf.*;

import org.apache.hadoop.io.*;

import org.apache.hadoop.mapreduce.*;

public class MatrixMapper extends Mapper <LongWritable, Text, Text, Text>

{

    public void map(LongWritable key, Text value, Context context) throws IOException,
    InterruptedException

    {

        Configuration conf = context.getConfiguration();

        int m = Integer.parseInt(conf.get("m"));

        int p = Integer.parseInt(conf.get("p"));

        String line = value.toString();

        String[] indicesAndValue = line.split(",");

        Text outputKey = new Text();

        Text outputValue = new Text();

        if (indicesAndValue[0].equals("M"))

        {

            for (int k = 0; k < p; k++)

            {

                outputKey.set(indicesAndValue[1] + "," + k);

                outputValue.set("M," + indicesAndValue[2] + "," + indicesAndValue[3]);

            }

        }

    }

}
```

```
context.write(outputKey, outputValue);

}

}

else

{

for (int i = 0; i < m; i++)

{

outputKey.set(i + ", " + indicesAndValue[2]);

outputValue.set("N, " + indicesAndValue[1] + ", " + indicesAndValue[3]);

context.write(outputKey, outputValue);

}

}

}

}
```

MatrixReducer.java

```
import java.io.IOException;

import java.util.*;

import org.apache.hadoop.io.*;

import org.apache.hadoop.mapreduce.*;

public class MatrixReducer extends Reducer<Text, Text, Text, Text>

{

public void reduce(Text key, Iterable<Text> values, Context

context) throws IOException, InterruptedException
```

```
{
String[] value;

HashMap<Integer, Float> hashA = new HashMap<Integer, Float>();

HashMap<Integer, Float> hashB = new HashMap<Integer, Float>();

for (Text val : values)
{
value = val.toString().split(",");

if (value[0].equals("M"))
{
hashA.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));
}

else
{
hashB.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));
}
}

int n = Integer.parseInt(context.getConfiguration().get("n")); float result
= 0.0f;

float a_ij; float b_jk;

for (int j = 0; j < n; j++)
{
a_ij = hashA.containsKey(j) ? hashA.get(j) : 0.0f;
b_jk = hashB.containsKey(j) ? hashB.get(j) : 0.0f;

result += a_ij * b_jk;
```

```

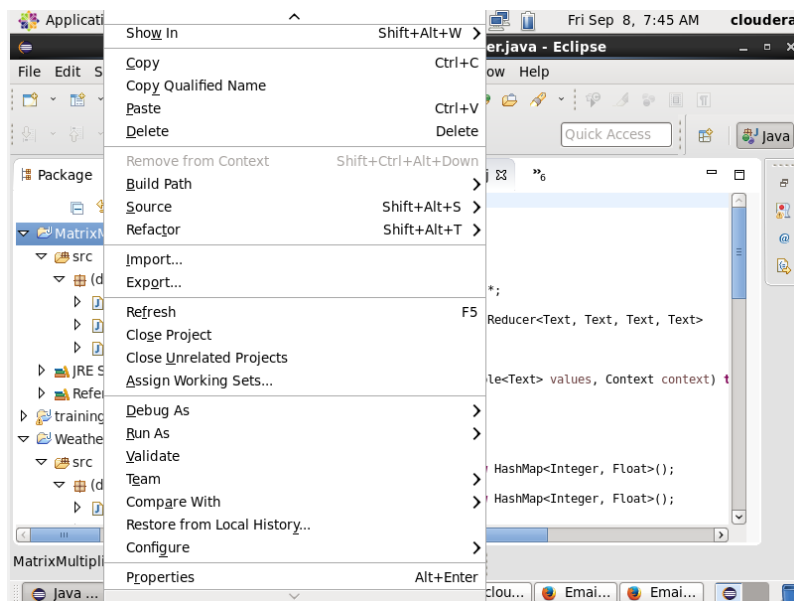
}

if (result != 0.0f)
{
context.write(null, new Text(key.toString() + "," + Float.toString(result)));
}

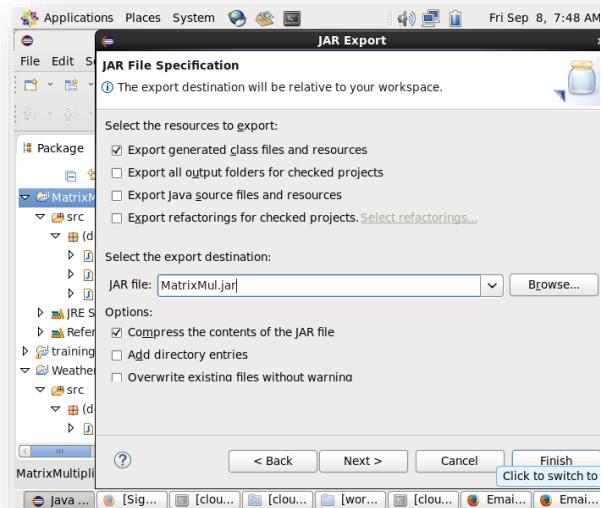
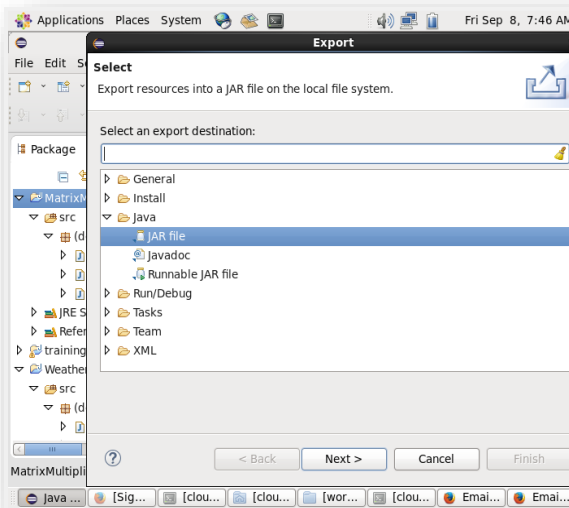
}

}

```



Right click on MatrixMultiplication->Export->java->jar file->JAR file:"**MatrixMul**"->Finish



OUTPUT:

In Terminal

```
cloudera@quickstart:~/workspace
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ cd workspace
[cloudera@quickstart workspace]$ cat > matrix.txt
M,0,0,1
M,0,1,2
M,1,0,3
M,1,1,4
N,0,0,5
N,0,1,6
N,1,0,7
N,1,1,8
^Z
[1]+  Stopped                  cat > matrix.txt
[cloudera@quickstart workspace]$
```



```
[cloudera@quickstart workspace]$ hadoop fs -put matrix.txt  
[cloudera@quickstart workspace]$
```

```
[cloudera@quickstart workspace]$ hadoop jar MatrixMul.jar MatrixDriver matrix.txt MatrixOutput
```

```
[cloudera@quickstart workspace]$ hadoop fs -cat MatrixOutput/part-r-00000  
0,0,19.0  
0,1,22.0  
1,0,43.0  
1,1,50.0
```

BDALAB